

# Red-Blue Pebbling Revisited: Near Optimal Parallel Matrix-Matrix Multiplication

Grzegorz Kwasniewski<sup>1</sup>, Marko Kabić<sup>2,3</sup>, Maciej Besta<sup>1</sup>,  
Joost VandeVondele<sup>2,3</sup>, Raffaele Solcà<sup>2,3</sup>, Torsten Hoefer<sup>1</sup>

<sup>1</sup>Department of Computer Science, ETH Zurich, <sup>2</sup>ETH Zurich, <sup>3</sup>Swiss National Supercomputing Centre (CSCS)

## ABSTRACT

We propose COSMA: a parallel matrix-matrix multiplication algorithm that is near communication-optimal for all combinations of matrix dimensions, processor counts, and memory sizes. The key idea behind COSMA is to derive an optimal (up to a factor of 0.03% for 10MB of fast memory) sequential schedule and then parallelize it, preserving I/O optimality. To achieve this, we use the red-blue pebble game to precisely model MMM dependencies and derive a constructive and tight sequential and parallel I/O lower bound proofs. Compared to 2D or 3D algorithms, which fix processor decomposition upfront and then map it to the matrix dimensions, it reduces communication volume by up to  $\sqrt{3}$  times. COSMA outperforms the established ScaLAPACK, CARMA, and CTF algorithms in all scenarios up to 12.8x (2.2x on average), achieving up to 88% of Piz Daint’s peak performance. Our work does not require any hand tuning and is maintained as an open source implementation.

## CCS CONCEPTS

- **Computing methodologies** → **Massively parallel algorithms**;
- **Mathematics of computing** → *Computations on matrices*;
- **Theory of computation** → **Distributed computing models**;
- **Communication complexity**; *Massively parallel algorithms*;

## ACM Reference Format:

Grzegorz Kwasniewski<sup>1</sup>, Marko Kabić<sup>2,3</sup>, Maciej Besta<sup>1</sup>, Joost VandeVondele<sup>2,3</sup>, Raffaele Solcà<sup>2,3</sup>, Torsten Hoefer<sup>1</sup> <sup>1</sup>Department of Computer Science, ETH Zurich, <sup>2</sup>ETH Zurich, <sup>3</sup>Swiss National Supercomputing Centre (CSCS). 2019. Red-Blue Pebbling Revisited: Near Optimal Parallel Matrix-Matrix Multiplication. In *The International Conference for High Performance Computing, Networking, Storage, and Analysis (SC ’19)*, November 17–22, 2019, Denver, CO, USA. ACM, New York, NY, USA, 13 pages. <https://doi.org/10.1145/3295500.3356181>

## 1 INTRODUCTION

Matrix-matrix multiplication (MMM) is one of the most fundamental building blocks in scientific computing, used in linear algebra algorithms [13, 15, 41], machine learning [6], graph processing [4, 8, 18, 36, 43, 51], computational chemistry [21], and others. Thus, accelerating MMM routines is of great significance for many

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from [permissions@acm.org](mailto:permissions@acm.org).

SC ’19, November 17–22, 2019, Denver, CO, USA

© 2019 Copyright held by the owner/author(s). Publication rights licensed to the Association for Computing Machinery.

ACM ISBN 978-1-4503-6229-0/19/11...\$15.00

<https://doi.org/10.1145/3295500.3356181>




Figure 1: Percentage of peak flop/s across the experiments ranging from 109 to 18,432 cores achieved by COSMA and the state-of-the-art libraries (Sec. 9).




Figure 2: Illustrative evolution of MMM algorithms reaching the I/O lower bound.

domains. In this work, we focus on minimizing the amount of transferred data in MMM, both across the memory hierarchy (*vertical I/O*) and between processors (*horizontal I/O*, aka “communication”)<sup>1</sup>.

The path to I/O optimality of MMM algorithms is at least 50 years old. The first parallel MMM algorithm is by Cannon [10], which works for square matrices and square processor decompositions. Subsequent works [24, 25] generalized the MMM algorithm to rectangular matrices, different processor decompositions, and communication patterns. PUMMA [17] package generalized previous works to transposed matrices and different data layouts. SUMMA algorithm [55] further extended it by optimizing the communication, introducing pipelining and communication-computation overlap. This is now a state-of-the-art so-called 2D algorithm (it decomposes processors in a 2D grid) used e.g., in ScaLAPACK library [14].

Agarwal et al. [1] showed that in a presence of extra memory, one can do better and introduces a 3D processor decomposition. The 2.5D algorithm by Solomonik and Demmel [52] effectively interpolates between those two results, depending on the available memory. However, Demmel et al. showed that algorithms optimized for square matrices often perform poorly when matrix dimensions vary significantly [22]. Such matrices are common in many relevant areas, for example in machine learning [59, 60] or computational chemistry [44, 48]. They introduced CARMA [22], a recursive algorithm that achieves asymptotic lower bounds for all configurations of dimensions and memory sizes. This evolution for chosen steps is depicted symbolically in Figure 2.

<sup>1</sup>We also focus only on “classical” MMM algorithms which perform  $n^3$  multiplications and additions. We do not analyze Strassen-like routines [53], as in practice they are often slower [19].

	2D [55]	2.5D [52]	recursive [22]	COSMA (this work)
Input:	User-specified grid	Available memory	Available memory, matrix dimensions	Available memory, matrix dimensions
<b>Step 1</b>	Split $m$ and $n$	Split $m, n, k$	Split recursively the largest dimension	Find the optimal sequential schedule
<b>Step 2</b>	Map matrices to processor grid	Map matrices to processor grid	Map matrices to recursion tree	Map sequential domain to matrices
	🔧 Requires manual tuning	🏆 Optimal for $m = n$	🏆 Asymptotically optimal for all $m, n, k, p$	🏆 Optimal for all $m, n, k$
	🔊 Asymptotically more comm.	🔊 Inefficient for $m \ll n$ or $n \ll m$	🔊 Up to $\sqrt{3}$ times higher comm. cost	🏆 Optimal for all $p$
		🔊 Inefficient for some $p$	🔊 $p$ must be a power of 2	🏆 Best time-to-solution

**Table 1:** Intuitive comparison between the COSMA algorithm and the state-of-the-art 2D, 2.5D, and recursive decompositions.  $C = AB$ ,  $A \in \mathbb{R}^{m \times k}$ ,  $B \in \mathbb{R}^{k \times n}$

Unfortunately, we observe several limitations with state-of-the-art algorithms. ScaLAPACK [14] (an implementation of SUMMA) supports only the 2D decomposition, which is communication-inefficient in the presence of extra memory. Also, it requires a user to fine-tune parameters such as block sizes or a processor grid size. CARMA supports only scenarios when the number of processors is a power of two [22], a serious limitation, as the number of processors is usually determined by the available hardware resources. Cyclops Tensor Framework (CTF) [49] (an implementation of the 2.5D decomposition) can utilize any number of processors, but its decompositions may be far from optimal (§ 9). We also emphasize that *asymptotic complexity is an insufficient measure of practical performance*. We later (§ 6.2) identify that CARMA performs up to  $\sqrt{3}$  more communication. Our observations are summarized in Table 1. Their practical implications are shown in Figure 1, where we see that all existing algorithms perform poorly for some configurations.

In this work, we present COSMA (Communication Optimal S-partition-based Matrix multiplication Algorithm): an algorithm that takes a new approach to multiplying matrices and alleviates the issues above. COSMA is I/O optimal for *all combinations of parameters* (up to the factor of  $\sqrt{S}/(\sqrt{S} + 1 - 1)$ , where  $S$  is the size of the fast memory<sup>2</sup>). The driving idea is to develop a general method of deriving I/O optimal schedules by explicitly modeling data reuse in the red-blue pebble game. We then parallelize the sequential schedule, minimizing the I/O between processors, and derive an optimal domain decomposition. This is in contrast with the other discussed algorithms, which fix the processor grid upfront and then map it to a sequential schedule for each processor. We outline the algorithm in § 3. To prove its optimality, we first provide a new constructive proof of a sequential I/O lower bound (§ 5.1), then we derive the communication cost of parallelizing the sequential schedule (§ 6.2), and finally we construct an I/O optimal parallel schedule (§ 6.3). The detailed communication analysis of COSMA, 2D, 2.5D, and recursive decompositions is presented in Table 2. Our algorithm reduces the data movement volume by a factor of up to  $\sqrt{3} \approx 1.73x$  compared to the asymptotically optimal recursive decomposition and up to  $\max\{m, n, k\}$  times compared to the 2D algorithms, like Cannon's [39] or SUMMA [55].

Our implementation enables transparent integration with the ScaLAPACK data format [16] and delivers near-optimal computation throughput. We later (§ 7) show that the schedule naturally expresses communication-computation overlap, enabling even higher speedups using Remote Direct Memory Access (RDMA). Finally, our I/O-optimal approach is generalizable to other linear algebra kernels. We provide the following contributions:

- We propose COSMA: a distributed MMM algorithm that is nearly-optimal (up to the factor of  $\sqrt{S}/(\sqrt{S} + 1 - 1)$ ) for *any combination of input parameters* (§ 3).
- Based on the red-blue pebble game abstraction [34], we provide a new method of deriving I/O lower bounds (Lemma 2), which may be used to generate optimal schedules (§ 4).
- Using Lemma 2, we provide a new constructive proof of the sequential MMM I/O lower bound. The proof delivers constant factors tight up to  $\sqrt{S}/(\sqrt{S} + 1 - 1)$  (§ 5).
- We extend the sequential proof to parallel machines and provide I/O optimal parallel MMM schedule (§ 6.3).
- We reduce memory footprint for communication buffers and guarantee minimal local data reshuffling by using a blocked data layout (§ 7.6) and a static buffer pre-allocation (§ 7.5), providing compatibility with the ScaLAPACK format.
- We evaluate the performance of COSMA, ScaLAPACK, CARMA, and CTF on the CSCS Piz Daint supercomputer for an extensive selection of problem dimensions, memory sizes, and numbers of processors, showing significant I/O reduction and the speedup of up to 8.3 times over the second-fastest algorithm (§ 9).

## 2 BACKGROUND

We first describe our machine model (§ 2.1) and computation model (§ 2.2). We then define our optimization goal: *the I/O cost* (§ 2.3).

### 2.1 Machine Model

We model a parallel machine with  $p$  processors, each with local memory of size  $S$  words. A processor can send and receive from any other processor up to  $S$  words at a time. To perform any computation, all operands must reside in processor's local memory. If shared memory is present, then it is assumed that it has infinite capacity. A cost of transferring a word from the shared to the local memory is equal to the cost of transfer between two local memories.

### 2.2 Computation Model

We now briefly specify a model of *general* computation; we use this model to derive the theoretical I/O cost in both the sequential and parallel setting. An execution of an algorithm is modeled with the *computational directed acyclic graph* (CDAG)  $G = (V, E)$  [11, 28, 46]. A vertex  $v \in V$  represents one elementary operation in the given computation. An edge  $(u, v) \in E$  indicates that an operation  $v$  depends on the result of  $u$ . A set of all immediate predecessors (or successors) of a vertex are its *parents* (or *children*). Two selected subsets  $I, O \subset V$  are *inputs* and *outputs*, that is, sets of vertices that have no parents (or no children, respectively).

#### Red-Blue Pebble Game

Hong and Kung's red-blue pebble game [34] models an execution of an algorithm in a two-level memory structure with a small-and-fast as well as large-and-slow memory. A red (or a blue) pebble placed

<sup>2</sup>Throughout this paper we use the original notation from Hong and Kung to denote the memory size  $S$ . In literature, it is also common to use the symbol  $M$  [2, 3, 33].

on a vertex of a CDAG denotes that the result of the corresponding elementary computation is inside the fast (or slow) memory. In the initial (or terminal) configuration, only inputs (or outputs) of the CDAG have blue pebbles. There can be at most  $S$  red pebbles used at any given time. A *complete CDAG calculation* is a sequence of moves that lead from the initial to the terminal configuration. One is allowed to: place a red pebble on any vertex with a blue pebble (load), place a blue pebble on any vertex with a red pebble (store), place a red pebble on a vertex whose parents all have red pebbles (compute), remove any pebble, red or blue, from any vertex (free memory). An *I/O optimal* complete CDAG calculation corresponds to a sequence of moves (called *pebbling* of a graph) which minimizes loads and stores. In the MMM context, it is an order in which all  $n^3$  multiplications are performed.

### 2.3 Optimization Goals

Throughout this paper we focus on the *input/output (I/O) cost* of an algorithm. The I/O cost  $Q$  is the total number of words transferred during the execution of a schedule. On a sequential or shared memory machine equipped with small-and-fast and slow-and-big memories, these transfers are load and store operations from and to the slow memory (also called the *vertical I/O*). For a distributed machine with a limited memory per node, the transfers are communication operations between the nodes (also called the *horizontal I/O*). A schedule is *I/O optimal* if it minimizes the I/O cost among all schedules of a given CDAG. We also model a *latency cost*  $L$ , which is a maximum number of messages sent by any processor.

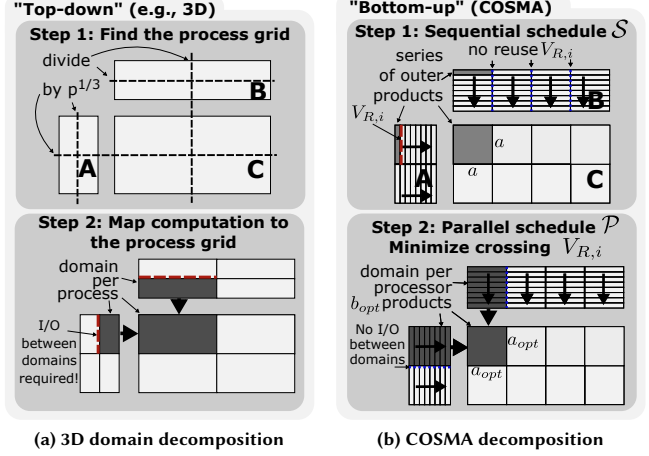
### 2.4 State-of-the-Art MMM Algorithms

Here we briefly describe strategies of the existing MMM algorithms. Throughout the whole paper, we consider matrix multiplication  $C = AB$ , where  $A \in \mathbb{R}^{m \times k}$ ,  $B \in \mathbb{R}^{k \times n}$ ,  $C \in \mathbb{R}^{m \times n}$ , where  $m$ ,  $n$ , and  $k$  are matrix dimensions. Furthermore, we assume that the size of each matrix element is one word, and that  $S < \min\{mn, mk, nk\}$ , that is, none of the matrices fits into single processor's fast memory.

We compare our algorithm with the 2D, 2.5D, and recursive decompositions (we select parameters for 2.5D to also cover 3D). We assume a square processor grid  $\lceil \sqrt{p}, \sqrt{p}, 1 \rceil$  for the 2D variant, analogously to Cannon's algorithm [10], and a cubic grid  $\lceil \sqrt[3]{p/c}, \sqrt[3]{p/c}, c \rceil$  for the 2.5D variant [52], where  $c$  is the amount of the "extra" memory  $c = pS/(mk + nk)$ . For the recursive decomposition, we assume that in each recursion level we split the largest dimension  $m$ ,  $n$ , or  $k$  in half, until the domain per processor fits into memory. The detailed complexity analysis of these decompositions is in Table 2. We note that ScaLAPACK or CTF can handle non-square decompositions, however they create different problems, as discussed in § 1. Moreover, in § 9 we compare their performance with COSMA and measure significant speedup in *all* scenarios.

### 3 COSMA: HIGH-LEVEL DESCRIPTION

COSMA decomposes processors by parallelizing the near optimal sequential schedule under constraints: (1) equal work distribution and (2) equal memory size per processor. Such a local sequential schedule is independent of matrix dimensions. Thus, intuitively, instead of dividing a global domain among  $p$  processors (the *top-down* approach), we start from deriving a near I/O optimal *sequential* schedule. We then parallelize it, minimizing the I/O and latency costs  $Q, L$  (the *bottom-up* approach); Figure 3 presents more details.



**Figure 3:** Domain decomposition using  $p = 8$  processors. In scenario (a), a straightforward 3D decomposition divides every dimension in  $p^{1/3} = 2$ . In scenario (b), COSMA starts by finding a near optimal sequential schedule and then parallelizes it minimizing crossing data reuse  $V_{R,i}$  (§ 5). The total communication volume is reduced by 17% compared to the former strategy.

COSMA is sketched in Algorithm 1. In Line 1 we derive a sequential schedule, which consists of series of  $a \times a$  outer products. (Figure 4 b). In Line 2, each processor is assigned to compute  $b$  of these products, forming a *local domain*  $\mathcal{D}$  (Figure 4 c), that is each  $\mathcal{D}$  contains  $a \times a \times b$  vertices (multiplications to be performed - the derivation of  $a$  and  $b$  is presented in § 6.3). In Line 3, we find a processor grid  $\mathcal{G}$  that evenly distributes this domain by the matrix dimensions  $m, n$ , and  $k$ . If the dimensions are not divisible by  $a$  or  $b$ , this function also evaluates new values of  $a_{opt}$  and  $b_{opt}$  by fitting the best matching decomposition, possibly not utilizing some processors (§ 7.1, Figure 4 d-f). The maximal number of idle processors is a tunable parameter  $\delta$ . In Line 5, we determine the initial decomposition of matrices  $A, B$ , and  $C$  to the submatrices  $A_l, B_l, C_l$  that are local for each processor. COSMA may handle any initial data layout, however, an optimal block-recursive one (§ 7.6) may be achieved in a preprocessing phase. In Line 6, we compute the size of the communication step, that is, how many of  $b_{opt}$  outer products assigned to each processor are computed in a single round, minimizing the latency (§ 6.3). In Line 7 we compute the number of sequential steps (Lines 8–11) in which every processor: (1) distributes and updates its local data  $A_l$  and  $B_l$  among the grid  $\mathcal{G}$  (Line 9), and (2) multiplies  $A_l$  and  $B_l$  (Line 10). Finally, the partial results  $C_l$  are reduced over  $\mathcal{G}$  (Line 12).

#### I/O Complexity of COSMA

Lines 2–7 require no communication (assuming that the parameters  $m, n, k, p, S$  are already distributed). The loop in Lines 8–11 executes  $\lceil 2ab/(S - a^2) \rceil$  times. In Line 9, each processor receives  $|A_l| + |B_l|$  elements. Sending the partial results in Line 12 adds  $a^2$  communicated elements. In § 6.3 we derive the optimal values for  $a$  and  $b$ , which yield a total of  $\min \left\{ S + 2 \cdot \frac{mnk}{p\sqrt{S}}, 3 \left( \frac{mnk}{p} \right)^{2/3} \right\}$  elements communicated.

### 4 ARBITRARY CDAGS: LOWER BOUNDS

We sketch key concepts and methods for deriving I/O lower bounds for general CDAGs. We extend the  $S$ -partition method and the related main lemma by Hong and Kung [34]. That lemma, however, fails to provide a tight I/O bound. In our generalization, the



**Algorithm 1** COSMA

---

**Input:** matrices  $A \in \mathbb{R}^{m \times k}$ ,  $B \in \mathbb{R}^{k \times n}$ ,  
number of processors:  $p$ , memory size:  $S$ , computation-I/O tradeoff ratio  $\rho$

**Output:** matrix  $C = AB \in \mathbb{R}^{m \times n}$

---

```

1:  $a \leftarrow \text{FindSeqSchedule}(S, m, n, k, p)$            ▶ sequential I/O optimality (§ 5)
2:  $b \leftarrow \text{ParallelizeSched}(a, m, n, k, p)$        ▶ parallel I/O optimality (§ 6)
3:  $(\mathcal{G}, a_{opt}, b_{opt}) \leftarrow \text{FitRanks}(m, n, k, a, b, p, \delta)$ 
4: for all  $p_i \in \{1 \dots p\}$  do in parallel
5:    $(A_i, B_i, C_i) \leftarrow \text{GetDataDecomp}(A, B, \mathcal{G}, p_i)$ 
6:    $s \leftarrow \left\lfloor \frac{S - a_{opt}}{2a_{opt}} \right\rfloor$            ▶ latency-minimizing size of a step (6.3)
7:    $t \leftarrow \left\lceil \frac{b_{opt}}{s} \right\rceil$            ▶ number of steps
8:   for  $j \in \{1 \dots t\}$  do
9:      $(A_i, B_i) \leftarrow \text{DistrData}(A_i, B_i, \mathcal{G}, j, p_i)$ 
10:     $C_i \leftarrow \text{Multiply}(A_i, B_i, j)$            ▶ compute locally
11:  end for
12:   $C \leftarrow \text{Reduce}(C_i, \mathcal{G})$            ▶ reduce the partial results
13: end for

```

---

**X-partition**, we let the constraint  $X$  to take any value (not only  $S$ , the fast memory size). Our key result here, Lemma 2, allows us to derive a constructive proof of a tight (for large  $S$ ) I/O lower bound for a sequential execution of the MMM CDAG (§ 5). This method is extended in Section 6 to establish a tight parallel I/O lower bound.

Due to space constraints, a formal definition of  $X$ -partition, proofs of all our lemmas, and a detailed discussion on their derivation and implications, are in an extended technical report<sup>3</sup>. Readers interested in practical concerns may proceed directly to Section 6.

**Reuse-based Lemma**

A proof of the original red-blue pebble lemma [34] is based on the fact that an optimal complete calculation (pebbling) which performs  $q$  I/O operations,  $hS \leq q < (h+1)S$  for some  $h$ , can be associated with an  $X$ -partition for  $X = 2S$ . Subsets (also called *subcomputations*)  $V_1, \dots, V_h \in V$  contain vertices that are red-pebbled *between every  $S$  I/O operations of this calculation* (i.e., between every  $(tS)$ -th and  $[(t+1)S]$ -th I/O operation,  $t \in \{1, \dots, h-1\}$ ). We refer the reader to the original paper for details of construction of sets  $V_i, i = 1 \dots h$ . A sketch of this construction is also in the technical report.

Before we present the lemmas, we need to introduce four more subsets of  $V$ , defined for each  $V_i$ .  $V_{R,i}$  (the *reuse set*) contains vertices that have red pebbles placed on them just before  $V_i$  begins (these vertices stay in the fast memory and are *reused* during  $V_i$ ).  $V_{B,i}$  (the *load set*) contains vertices that have blue pebbles placed on them just before  $V_i$  begins, and have red pebbles placed on them during  $V_i$  (for these vertices, we need to load data from the slow to the fast memory). By definition<sup>4</sup>,  $\text{Dom}(V_i) = V_{R,i} \cup V_{B,i}$ . We define similar subsets  $W_{B,i}$  and  $W_{R,i}$  for the minimum set<sup>5</sup>  $\text{Min}(V_i)$ .  $W_{B,i}$  (the *store set*) contains all vertices in  $V_i$  that have a blue pebble placed on them during  $V_i$  (i.e., the corresponding data must be stored in the slow memory).  $W_{R,i}$  (called the *cache set*) contains all vertices in  $V_i$  that have a red pebble at the end of  $V_i$  (the data resides in the fast memory after  $V_i$ ).

Denote an upper bound on  $|V_{R,i}|$  and  $|W_{B,i}|$  as  $R(S)$  ( $\forall i \max\{|V_{R,i}|, |W_{B,i}|\} \leq R(S) \leq S$ ). Further, denote a lower bound on  $|V_{B,i}|$  and  $|W_{R,i}|$  as  $T(S)$  ( $\forall i 0 \leq T(S) \leq \min\{|V_{B,i}|, |W_{R,i}|\}$ ). We now use  $R(S)$  and  $T(S)$  to tighten the bound on  $Q$ .

We now use the above definitions and observations to **generalize the result of Hong and Kung [34]**.

LEMMA 1. Denote  $H(X)$  as the minimum number of subcomputations in any valid  $X$ -partition of a CDAG  $G = (V, E)$ , for any  $X \geq S$ . The minimal number  $Q$  of I/O operations for any valid execution of a CDAG  $G = (V, E)$  is bounded by

$$Q \geq (X - R(S) + T(S)) \cdot (H(X) - 1) \quad (1)$$

where  $R(S)$  is the maximum reuse set size and  $T(S)$  is the minimum store set size. Moreover, we have

$$H(X) \geq \frac{|V|}{|V_{max}|} \quad (2)$$

where  $V_{max} = \arg \max_{V_i \in S(X)} |V_i|$  is the largest subset of vertices in the CDAG schedule  $S(X) = \{V_1, \dots, V_h\}$ .

From this lemma, we derive the following lemma that we use to prove a tight I/O lower bound for MMM (Theorem 1):

LEMMA 2. Define the number of computations performed by  $V_i$  for one loaded element as the computational intensity  $\rho_i = \frac{|V_i|}{X - |V_{R,i}| + |W_{B,i}|}$  of the subcomputation  $V_i$ . Denote  $\rho = \max_i(\rho_i) \leq \frac{|V_{max}|}{X - R(S) + T(S)}$  to be the maximal computational intensity. Then, the number of I/O operations  $Q$  is bounded by  $Q \geq |V|/\rho$ .

**5 NEAR-OPTIMAL SEQUENTIAL MMM**

In this section, we present our main theoretical contribution: a constructive proof of a tight I/O lower bound for classical matrix-matrix multiplication. In § 6, we extend it to the parallel setup (Theorem 2). This result is tight (up to diminishing factor  $\sqrt{S}/(\sqrt{S} + 1 - 1)$ ), and therefore may be seen as the last step in the long sequence of improved bounds. Hong and Kung [34] derived an asymptotic bound  $\Omega(n^3/\sqrt{S})$  for the sequential case. Irony et al. [33] extended the lower bound result to a parallel machine with  $p$  processes, each having a fast private memory of size  $S$ , proving the  $\frac{n^3}{4\sqrt{2}p\sqrt{S}} - S$  lower bound on the communication volume per process. Recently, Smith and van de Gein [47] proved a tight sequential lower bound (up to an additive term) of  $2mnk/\sqrt{S} - 2S$ . Our proof improves the additive term and extends it to a parallel schedule.

THEOREM 1 (SEQUENTIAL MATRIX MULTIPLICATION I/O LOWER BOUND). Any pebbling of MMM CDAG which multiplies matrices of sizes  $m \times k$  and  $k \times n$  by performing  $mnk$  multiplications requires a minimum number of  $\frac{2mnk}{\sqrt{S}} + mn$  I/O operations.

The proof of Theorem 1 requires Lemmas 3 and 4, which in turn, require a following definition:

**Greedy schedule**

A schedule  $\mathcal{S} = \{V_1, \dots, V_h\}$  is *greedy* if during every subcomputation  $V_r$  every vertex  $u$  that will hold a red pebble either has a child in  $V_r$  or belongs to  $V_r$ . Then, our first result is an I/O lower bound for such schedules:

LEMMA 3. Any greedy schedule that multiplies matrices of sizes  $m \times k$  and  $k \times n$  using  $mnk$  multiplications requires a minimum number of  $\frac{2mnk}{\sqrt{S}} + mn$  I/O operations.

<sup>3</sup>Available at <https://arxiv.org/abs/1908.09606>

<sup>4</sup>A dominator set  $\text{Dom}(V_i)$  is a set of vertices in  $V$ , such that every path from any input of a CDAG to any vertex in  $V_i$  must contain at least one vertex in  $\text{Dom}(V_i)$ .

<sup>5</sup>The minimum set  $\text{Min}(V_i)$  is a set of all vertices in  $V_i$  with no children in  $V_i$ .

*Intuition: Restricting the analysis to greedy schedules provides explicit information of a state of memory (sets  $V_r$ ,  $V_{R,r}$ ,  $W_{B,r}$ ), and to a corresponding CDAG pebbling. Additional constraints (§ 5.1) guarantee feasibility of a derived schedule (and therefore, lower bound tightness).*

### 5.1 Near-Optimal Greedy Schedule

In the technical report, it is proven that an optimal greedy schedule is composed of  $\frac{mnk}{R(S)}$  outer product calculations, while loading  $\sqrt{R(S)}$  elements of each of matrices  $A$  and  $B$ . While the lower bound is achieved for  $R(S) = S$ , such a schedule is infeasible, as at least some additional red pebbles, except the ones placed on the reuse set  $V_{R,r}$ , have to be placed on  $2\sqrt{R(S)}$  vertices of  $A$  and  $B$ .

A direct way to obtain a feasible greedy schedule is to set  $X = S$ , ensuring that the dominator set can fit into the memory. Then each subcomputation is an outer-product of column-vector of matrix  $A$  and row-vector of  $B$ , both holding  $\sqrt{S} + 1 - 1$  values. Such a schedule performs  $\frac{2mnk}{\sqrt{S}+1-1} + mn$  I/O operations, a factor of  $\frac{\sqrt{S}}{\sqrt{S}+1-1}$  more than a lower bound, which quickly approach 1 for large  $S$ . Listing 1 provides a pseudocode of this algorithm, which is a well-known rank-1 update formulation of MMM.

```

1 for  $i_1 = 1 : \lceil m/a \rceil$ 
2   for  $j_1 = 1 : \lceil n/a \rceil$ 
3     for  $r = 1 : k$ 
4       for  $i_2 = i_1 \cdot T : \min((i_1 + 1) \cdot a, m)$ 
5         for  $j_2 = j_1 \cdot T : \min((j_1 + 1) \cdot a, n)$ 
6            $C(i_2, j_2) = C(i_2, j_2) + A(i_2, r) \cdot B(r, j_2)$ 

```

Listing 1: Pseudocode of near optimal sequential MMM,  $a = \sqrt{S} + 1 - 1$

### 5.2 Greedy vs Non-greedy Schedules

Greedy schedules provide means to directly derive near-optimal schedules from the corresponding lower bounds, as shown in § 5.1. However, restricting analysis just to them may not provide correct lower bounds for general CDAGs. In this section, a different approach is used to prove the lower bound for non-greedy schedules - however, without a direct way to retrieve the corresponding schedule. These two bounds match, proving that the derived schedule (Listing 1) is near-optimal.

LEMMA 4. *Any non-greedy schedule computing classical matrix multiplication performs at least  $\frac{2mnk}{\sqrt{S}} + mn$  I/O operations.*

#### Proof of Theorem 1:

Lemma 3 establishes that the I/O lower bound for any greedy schedule is  $Q = 2mnk/\sqrt{S} + mn$ . Lemma 4 establishes that no other schedule can perform less I/O operations.  $\square$

## 6 OPTIMAL PARALLEL MMM

We now derive the schedule of COSMA from the results from § 5.1. The key notion is the data reuse, that determines not only the sequential execution, as discussed in § 4, but also the parallel scheduling. Specifically, if the data reuse set spans across multiple local domains, then this set has to be communicated between these domains, increasing the I/O cost (Figure 3). We first introduce a formalism required to parallelize the sequential schedule (§ 6.1). In § 6.2, we generalize parallelization strategies used by the 2D,

2.5D, and recursive decompositions, deriving their communication cost and showing that none of them is optimal in the whole range of parameters. We finally derive the optimal decomposition (*FindOptimalDomain* function in Algorithm 1) by expressing it as an optimization problem (§ 6.3), and analyzing its I/O and latency cost. The remaining steps in Algorithm 1: *FitRanks*, *GetDataDecomp*, as well as *DistrData* and *Reduce* are discussed in § 7.1, § 7.6, and § 7.2, respectively. For a distributed machine, we assume that all matrices fit into collective memories of all processors:  $pS \geq mn + mk + nk$ . For a shared memory setting, we assume that all inputs start in a common slow memory.

### 6.1 Sequential and Parallel Schedules

We now describe how a parallel schedule is formed from a sequential one. The sequential schedule  $\mathcal{S}$  partitions the CDAG  $G = (V, E)$  into  $H(S)$  subcomputations  $V_i$ . The parallel schedule  $\mathcal{P}$  divides  $\mathcal{S}$  among  $p$  processors:  $\mathcal{P} = \{\mathcal{D}_1, \dots, \mathcal{D}_p\}$ ,  $\bigcup_{j=1}^p \mathcal{D}_j = \mathcal{S}$ . The set  $\mathcal{D}_j$  of all  $V_k$  assigned to processor  $j$  forms a *local domain* of  $j$  (Fig. 4c).

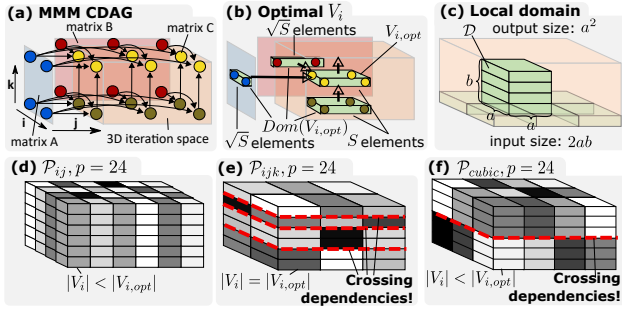
If two local domains  $\mathcal{D}_k$  and  $\mathcal{D}_l$  are dependent, that is,  $\exists u, \exists v : u \in \mathcal{D}_k \wedge v \in \mathcal{D}_l \wedge (u, v) \in E$ , then  $u$  has to be *communicated* from processor  $k$  to  $l$ . The total number of vertices communicated between all processors is the *I/O cost*  $Q$  of schedule  $\mathcal{P}$ . We say that the parallel schedule  $\mathcal{P}_{opt}$  is *communication-optimal* if  $Q(\mathcal{P}_{opt})$  is minimal among all possible parallel schedules.

The vertices of MMM CDAG may be arranged in an  $[m \times n \times k]$  3D grid called an *iteration space* [58]. The orthonormal vectors  $\mathbf{i}, \mathbf{j}, \mathbf{k}$  correspond to the loops in Lines 1-3 in Listing 1 (Figure 3a). We call a schedule  $\mathcal{P}$  *parallelized in dimension  $\mathbf{d}$*  if we “cut” the CDAG along dimension  $\mathbf{d}$ . More formally, each local domain  $\mathcal{D}_j$ ,  $j = 1 \dots p$  is a grid of size either  $[m/p, n, k]$ ,  $[m, n/p, k]$ , or  $[m, n, k/p]$ . The schedule may also be parallelized in two dimensions ( $\mathbf{d}_1 \mathbf{d}_2$ ) or three dimensions ( $\mathbf{d}_1 \mathbf{d}_2 \mathbf{d}_3$ ) with a local domain size  $[m/p_m, n/p_n, k/p_k]$  for some  $p_m, p_n, p_k$ , such that  $p_m p_n p_k = p$ . We call  $\mathcal{G} = [p_m, p_n, p_k]$  the *processor grid* of a schedule. E.g., Cannon’s algorithm is parallelized in dimensions  $\mathbf{ij}$ , with the processor grid  $[\sqrt{p}, \sqrt{p}, 1]$ . COSMA, on the other hand, may use any of the possible parallelizations, depending on the problem parameters.

### 6.2 Parallelization Strategies for MMM

The sequential schedule  $\mathcal{S}$  (§ 5) consists of  $mnk/S$  elementary outer product calculations, arranged in  $\sqrt{S} \times \sqrt{S} \times k$  “blocks” (Figure 4). The number  $p_1 = mn/S$  of dependency-free subcomputations  $V_i$  (i.e., having no parents except for input vertices) in  $\mathcal{S}$  determines the maximum degree of parallelism of  $\mathcal{P}_{opt}$  for which no reuse set  $V_{R,i}$  crosses two local domains  $\mathcal{D}_j, \mathcal{D}_k$ . The optimal schedule is parallelized in dimensions  $\mathbf{ij}$ . There is no communication between the domains (except for inputs and outputs), and all I/O operations are performed inside each  $\mathcal{D}_j$  following the sequential schedule. Each processor is assigned to  $p_1/p$  local domains  $\mathcal{D}_j$  of size  $[\sqrt{S}, \sqrt{S}, k]$ , each of which requires  $2\sqrt{S}k + S$  I/O operations (Theorem 1), giving a total of  $Q = 2mnk/(p\sqrt{S}) + mn/p$  I/O operations per processor.

When  $p > p_1$ , the size of local domains  $|\mathcal{D}_j|$  is smaller than  $\sqrt{S} \times \sqrt{S} \times k$ . Then, the schedule has to either be parallelized in dimension  $\mathbf{k}$ , or has to reduce the size of the domain in  $\mathbf{ij}$  plane. The former option creates dependencies between the local domains, which results in additional communication (Figure 4e). The latter



**Figure 4:** (a) An MMM CDAG as a 3D grid (iteration space). Each vertex in it (except for the vertices in the bottom layer) has three parents - blue (matrix A), red (matrix B), and yellow (partial result of matrix C) and one yellow child (except for vertices in the top layer). (b) A union of inputs of all vertices in  $V_i$  form the dominator set  $Dom(V_i)$  (two blue, two red and four dark yellow). Using approximation  $\sqrt{S} + 1 - 1 \approx \sqrt{S}$ , we have  $|Dom(V_{i,opt})| = S$ . (c) A local domain  $\mathcal{D}$  consists of  $b$  subcomputations  $V_i$ , each of a dominator size  $|Dom(V_i)| = a^2 + 2a$ . (d-f) Different parallelization schemes of near optimal sequential MMM for  $p = 24 > p_1 = 6$ .

does not utilize the whole available memory, making the sequential schedule not I/O optimal and decreasing the computational intensity  $\rho$  (Figure 4d). We now analyze three possible parallelization strategies (Figure 4) which generalize 2D, 2.5D, and recursive decomposition strategies; see Table 2 for details.

**Schedule  $\mathcal{P}_{ij}$**  The schedule is parallelized in dimensions  $i$  and  $j$ . The processor grid is  $\mathcal{G}_{ij} = [\frac{m}{a}, \frac{n}{a}, 1]$ , where  $a = \sqrt{\frac{mn}{p}}$ . Because all dependencies are parallel to dimension  $k$ , there are no dependencies between  $\mathcal{D}_j$  except for the inputs and the outputs. Because  $a < \sqrt{S}$ , the corresponding sequential schedule has a reduced computational intensity  $\rho_{ij} < \sqrt{S}/2$ .

**Schedule  $\mathcal{P}_{ijk}$**  The schedule is parallelized in all dimensions. The processor grid is  $\mathcal{G}_{ijk} = [\frac{m}{\sqrt{S}}, \frac{n}{\sqrt{S}}, \frac{k}{pS}]$ . The computational intensity  $\rho_{ijk} = \sqrt{S}/2$  is optimal. The parallelization in  $k$  dimension creates dependencies between local domains, requiring communication and increasing the I/O cost.

**Schedule  $\mathcal{P}_{cubic}$**  The schedule is parallelized in all dimensions. The grid is  $[\frac{m}{a_c}, \frac{n}{a_c}, \frac{k}{a_c}]$ , where  $a_c = \min\left\{\left(\frac{mnk}{p}\right)^{1/3}, \sqrt{\frac{S}{3}}\right\}$ . Because  $a_c < \sqrt{S}$ , the corresponding computational intensity  $\rho_{cubic} < \sqrt{S}/2$  is not optimal. The parallelization in  $k$  dimension creates dependencies between local domains, increasing communication.

#### Schedules of the State-of-the-Art Decompositions

If  $m = n$ , the  $\mathcal{P}_{ij}$  scheme is reduced to the classical 2D decomposition (e.g., Cannon's algorithm [10]), and  $\mathcal{P}_{ijk}$  is reduced to the 2.5D decomposition [52]. CARMA [22] asymptotically reaches the  $\mathcal{P}_{cubic}$  scheme, guaranteeing that the longest dimension of a local cuboidal domain is at most two times larger than the smallest one. We present a detailed complexity analysis comparison for all algorithms in Table 2.

### 6.3 I/O Optimal Parallel Schedule

Observe that none of those schedules is optimal in the whole range of parameters. As discussed in § 5, in sequential scheduling, intermediate results of  $C$  are not stored to the memory: they are consumed (reused) immediately by the next sequential step. Only the final result of  $C$  in the local domain is sent. Therefore, the optimal parallel

schedule  $\mathcal{P}_{opt}$  minimizes the communication, that is, sum of the inputs' sizes plus the output size, under the sequential I/O constraint on subcomputations  $\forall V_i \in \mathcal{D}_j \in \mathcal{P}_{opt} |Dom(V_i)| \leq S \wedge |Min(V_i)| \leq S$ .

The local domain  $\mathcal{D}_j$  is a grid of size  $[a \times a \times b]$ , containing  $b$  outer products of vectors of length  $a$ . The optimization problem of finding  $\mathcal{P}_{opt}$  using the computational intensity (Lemma 2) is formulated as follows:

$$\begin{aligned} \text{maximize } \rho &= \frac{a^2 b}{ab + ab + a^2} \\ \text{subject to:} \end{aligned} \quad (3)$$

$$a^2 \leq S \text{ (the I/O constraint)}$$

$$a^2 b = \frac{mnk}{p} \text{ (the load balance constraint)}$$

$$pS \geq mn + mk + nk \text{ (matrices must fit into memory)}$$

The I/O constraint  $a^2 \leq S$  is binding (changes to equality) for  $p \geq \frac{mnk}{S^{3/2}}$ . Therefore, the solution to this problem is:

$$a = \min\left\{\sqrt{S}, \left(\frac{mnk}{p}\right)^{1/3}\right\}, \quad b = \max\left\{\frac{mnk}{pS}, \left(\frac{mnk}{p}\right)^{1/3}\right\} \quad (4)$$

The I/O complexity of this schedule is:

$$Q \geq \frac{a^2 b}{\rho} = \min\left\{\frac{2mnk}{p\sqrt{S}} + S, 3\left(\frac{mnk}{p}\right)^{2/3}\right\} \quad (5)$$

This can be intuitively interpreted geometrically as follows: if we imagine the optimal local domain "growing" with the decreasing number of processors, then it stays cubic as long as it is still "small enough" (its side is smaller than  $\sqrt{S}$ ). After that point, its face in the  $ij$  plane stays constant  $\sqrt{S} \times \sqrt{S}$  and it "grows" only in the  $k$  dimension. This schedule effectively switches from  $\mathcal{P}_{ijk}$  to  $\mathcal{P}_{cubic}$  once there is enough memory ( $S \geq (mnk/p)^{2/3}$ ).

**THEOREM 2.** The I/O complexity of a classic Matrix Multiplication algorithm executed on  $p$  processors, each of local memory size  $S \geq \frac{mn+mk+nk}{p}$  is

$$Q \geq \min\left\{\frac{2mnk}{p\sqrt{S}} + S, 3\left(\frac{mnk}{p}\right)^{2/3}\right\}$$

**PROOF.** The theorem is a direct consequence of Lemma 1 and the computational intensity (Lemma 2). The load balance constraint enforces a size of each local domain  $|\mathcal{D}_j| = mnk/p$ . The I/O cost is then bounded by  $|\mathcal{D}_j|/\rho$ . Schedule  $\mathcal{P}_{opt}$  maximizes  $\rho$  by the formulation of the optimization problem (Equation 3).  $\square$

#### I/O-Latency Trade-off

As showed in this section, the local domain  $\mathcal{D}$  of the near optimal schedule  $\mathcal{P}$  is a grid of size  $[a \times a \times b]$ , where  $a, b$  are given by Equation (4). The corresponding sequential schedule  $\mathcal{S}$  is a sequence of  $b$  outer products of vectors of length  $a$ . Denote the size of the communicated inputs in each step by  $I_{step} = 2a$ . This corresponds to  $b$  steps of communication (the latency cost is  $L = b$ ).

The number of steps (latency) is equal to the total communication volume of  $\mathcal{D}$  divided by the volume per step  $L = Q/I_{step}$ . To reduce the latency, one either has to decrease  $Q$  or increase  $I_{step}$ , under the memory constraint that  $I_{step} + a^2 \leq S$  (otherwise we cannot fit both the inputs and the outputs in the memory). Express  $I_{step} = a \cdot h$ ,

Decomposition Parallel schedule $\mathcal{P}$	2D [55] $\mathcal{P}_{ij}$ for $m = n$	2.5D [52] $\mathcal{P}_{ijk}$ for $m = n$	recursive [22] $\mathcal{P}_{cubic}$	COSMA (this paper) $\mathcal{P}_{opt}$
grid $[p_m \times p_n \times p_k]$	$[\sqrt{p} \times \sqrt{p} \times 1]$	$[\sqrt{p/c} \times \sqrt{p/c} \times c]; c = \frac{pS}{mk+nk}$	$[2^{a_1} \times 2^{a_2} \times 2^{a_3}]; a_1 + a_2 + a_3 = \log_2(p)$	$[\frac{m}{a} \times \frac{n}{a} \times \frac{k}{b}]; a, b : \text{Equation 4}$
domain size	$[\frac{m}{\sqrt{p}} \times \frac{n}{\sqrt{p}} \times k]$	$[\frac{m}{\sqrt{p/c}} \times \frac{n}{\sqrt{p/c}} \times \frac{k}{c}]$	$[\frac{m}{2^{a_1}} \times \frac{n}{2^{a_2}} \times \frac{k}{2^{a_3}}]$	$[a \times a \times b]$
<b>"General case":</b>				
I/O cost $Q$	$\frac{k}{\sqrt{p}}(m+n) + \frac{mn}{p}$	$\frac{(k(m+n))^{3/2}}{p\sqrt{S}} + \frac{mnS}{k(m+n)}$	$2 \min \left\{ \sqrt{3} \frac{mnk}{p\sqrt{S}}, \left( \frac{mnk}{p} \right)^{2/3} \right\} + \left( \frac{mnk}{p} \right)^{2/3}$	$\min \left\{ \frac{2mnk}{p\sqrt{S}} + S, 3 \left( \frac{mnk}{p} \right)^{2/3} \right\}$
latency cost $L$	$2k \log_2(\sqrt{p})$	$\frac{(k(m+n))^{5/2}}{pS^{3/2}(k+m+k-mn)} + 3 \log_2 \left( \frac{pS}{mk+nk} \right)$	$(3^{3/2} mnk) / (pS^{3/2}) + 3 \log_2(p)$	$\frac{2ab}{S-a^2} \log_2 \left( \frac{mn}{a^2} \right)$
<b>Square matrices, "limited memory": <math>m = n = k, S = 2n^2/p, p = 2^{3n}</math></b>				
I/O cost $Q$	$2n^2(\sqrt{p}+1)/p$	$2n^2(\sqrt{p}+1)/p$	$2n^2 \left( \sqrt{3/2p} + 1/2p^{2/3} \right)$	$2n^2(\sqrt{p}+1)/p$
latency cost $L$	$2k \log_2(\sqrt{p})$	$\sqrt{p}$	$(\frac{3}{2})^{3/2} \sqrt{p} \log_2(p)$	$\sqrt{p} \log_2(p)$
<b>"Tall" matrices, "extra" memory available: <math>m = n = \sqrt{p}, k = p^{3/2}/4, S = 2nk/p^{2/3}, p = 2^{3n+1}</math></b>				
I/O cost	$p^{3/2}/2$	$p^{4/3}/2 + p^{1/3}$	$3p/4$	$p \left( 3 - 2^{1/3} \right) / 2^{4/3} \approx 0.69p$
latency cost $L$	$p^{3/2} \log_2(\sqrt{p})/4$	1	1	1

**Table 2:** The comparison of complexities of 2D, 2.5D, recursive, and COSMA algorithms. The 3D decomposition is a special case of 2.5D, and can be obtained by instantiating  $c = p^{1/3}$  in the 2.5D case. In addition to the general analysis, we show two special cases. If the matrices are square and there is no extra memory available, 2D, 2.5D and COSMA achieves tight communication lower bound  $2n^2/\sqrt{p}$ , whereas CARMA performs  $\sqrt{3}$  times more communication. If one dimension is much larger than the others and there is extra memory available, 2D, 2.5D and CARMA decompositions perform  $O(p^{1/2})$ ,  $O(p^{1/3})$ , and 8% more communication than COSMA, respectively. For simplicity, we assume that parameters are chosen such that all divisions have integer results.


where  $h$  is the number of sequential subcomputations  $V_i$  we merge in one communication. We can express the I/O-latency trade-off:

$$\begin{aligned}
 & \min(Q, L) \\
 & \text{subject to:} \\
 & Q = 2ab + a^2, L = \frac{b}{h} \\
 & a^2 + 2ah \leq S \text{ (I/O constraint)} \\
 & a^2b = \frac{mnk}{p} \text{ (load balance constraint)}
 \end{aligned}$$

Solving this problem, we have  $Q = \frac{2mnk}{pa} + a^2$  and  $L = \frac{2mnk}{pa(S-a^2)}$ , where  $a \leq \sqrt{S}$ . Increasing  $a$  we *reduce* the I/O cost  $Q$  and *increase* the latency cost  $L$ . For minimal value of  $Q$  (Theorem 2),  $L = \left\lceil \frac{2ab}{S-a^2} \right\rceil$ , where  $a = \min\{\sqrt{S}, (mnk/p)^{1/3}\}$  and  $b = \max\{\frac{mnk}{pS}, (mnk/p)^{1/3}\}$ . Based on our experiments, we observe that the I/O cost is vastly greater than the latency cost, therefore our schedule by default minimizes  $Q$  and uses extra memory (if any) to reduce  $L$ .

## 7 IMPLEMENTATION

We now present implementation optimizations that further increase the performance of COSMA on top of the speedup due to our near I/O optimal schedule. The algorithm is designed to facilitate the overlap of communication and computation § 7.3. For this, to leverage the RDMA mechanisms of current high-speed network interfaces, we use the MPI one-sided interface § 7.4. In addition, our implementation also offers alternative efficient two-sided communication back end that uses MPI collectives. We also use a blocked data layout § 7.6, a grid-fitting technique § 7.1, and an optimized binary broadcast tree using static information about the communication pattern (§ 7.2) together with the buffer swapping (§ 7.5). For the local matrix operations, we use BLAS routines for highest performance. Our code is publicly available at <https://github.com/eth-cscs/COSMA>.



**Figure 5:** Processor decomposition for square matrices and 65 processors. (a) To utilize all resources, the local domain is drastically stretched. (b) Dropping one processor results in a symmetric grid which increases the computation per processor by 1.5%, but reduces the communication by 36%.

### 7.1 Processor Grid Optimization

Throughout the paper, we assume all operations required to assess the decomposition (divisions, roots) result in natural numbers. We note that in practice it is rarely the case, as the parameters usually emerge from external constraints, like a specification of a performed calculation or hardware resources (§ 8). If matrix dimensions are not divisible by the local domain sizes  $a, b$  (Equation 4), then a straightforward option is to use the floor function, not utilizing the "boundary" processors whose local domains do not fit entirely in the iteration space, which result in more computation per processor. The other option is to find factors of  $p$  and then construct the processor grid by matching the largest factors with largest matrix dimensions. However, if the factors of  $p$  do not match  $m, n$ , and  $k$ , this may result in a suboptimal decomposition. Our algorithm allows to not utilize some processors (increasing the computation volume per processor) to optimize the grid, which reduces the communication volume. Figure 5 illustrates the comparison between these options. We balance this communication-computation trade-off by "stretching" the local domain size derived in § 6.3 to fit the global domain by adjusting its width, height, and length. The range of this tuning (how many processors we drop to reduce communication) depends on the hardware specification of the machine (peak flop/s, memory and network bandwidth). For our experiments on the Piz Daint machine, we chose the maximal number of unutilized cores to be 3%, accounting for up to 2.4 times speedup for the square matrices using 2,198 cores (§ 9).



## 7.2 Enhanced Communication Pattern

As shown in Algorithm 1, COSMA by default executes in  $t = \frac{2ab}{S-a^2}$  rounds. In each round, each processor receives  $s = ab/t = (S-a^2)/2$  elements of  $A$  and  $B$ . Thus, the input matrices are broadcast among the  $i$  and  $j$  dimensions of the processor grid. After the last round, the partial results of  $C$  are reduced among the  $k$  dimension. The communication pattern is therefore similar to ScaLAPACK or CTF.

To accelerate the collective communication, we implement our own binary broadcast tree, taking advantage of the known data layout, processor grid, and communication pattern. Knowing the initial data layout § 7.6 and the processor grid § 7.1, we craft the binary reduction tree in all three dimensions  $i$ ,  $j$ , and  $k$  such that the distance in the grid between communicating processors is minimized. Our implementation outperforms the standard MPI broadcast from the Cray-MPICH 3.1 library by approximately 10%.

## 7.3 Communication–Computation Overlap

The sequential rounds of the algorithm  $t_i = 1, \dots, t$ , naturally express communication–computation overlap. Using double buffering, at each round  $t_i$  we issue an asynchronous communication (using either MPI\_Get or MPI\_Isend / MPI\_Irecv § 7.4) of the data required at round  $t_{i+1}$ , while locally processing the data received in a previous round. We note that, by the construction of the local domains  $\mathcal{D}_j$  § 6.3, the extra memory required for double buffering is rarely an issue. If we are constrained by the available memory, then the space required to hold the partial results of  $C$ , which is  $a^2$ , is much larger than the size of the receive buffers  $s = (S-a^2)/2$ . If not, then there is extra memory available for the buffering.

### Number of rounds

The minimum number of rounds, and therefore latency, is  $t = \frac{2ab}{S-a^2}$  (§ 6.3). However, to exploit more overlap, we can increase the number of rounds  $t_2 > t$ . In this way, in one round we communicate less data  $s_2 = ab/t_2 < s$ , allowing the first round of computation to start earlier.

## 7.4 One-Sided vs Two-Sided Communication

To reduce the latency [27] we implemented communication using MPI RMA [32]. This interface utilizes the underlying features of Remote Direct Memory Access (RDMA) mechanism, bypassing the OS on the sender side and providing zero-copy communication: data sent is not buffered in a temporary address, instead, it is written directly to its location.

All communication windows are pre-allocated using MPI\_Win\_allocate with the size of maximum message in the broadcast tree  $2^{s-1}D$  (§ 7.2). Communication in each step is performed using the MPI\_Get and MPI\_Accumulate routines.

For compatibility reasons, as well as for the performance comparison, we also implemented a communication back-end using MPI two-sided (the message passing abstraction).

## 7.5 Communication Buffer Optimization

The binary broadcast tree pattern is a generalization of the recursive structure of CARMA. However, CARMA in each recursive step dynamically allocates new buffers of the increasing size to match the message sizes  $2^{s-1}D$ , causing an additional runtime overhead.

To alleviate this problem, we pre-allocate initial, send, and receive buffers for each of matrices  $A$ ,  $B$ , and  $C$  of the maximum size

of the message  $ab/t$ , where  $t = \frac{2ab}{S-a^2}$  is the number of steps in COSMA (Algorithm 1). Then, in each level  $s$  of the communication tree, we move the pointer in the receive buffer by  $2^{s-1}D$  elements.

## 7.6 Blocked Data Layout

COSMA’s schedule induces the optimal initial data layout, since for each  $\mathcal{D}_j$  it determines its dominator set  $Dom(\mathcal{D}_j)$ , that is, elements accessed by processor  $j$ . Denote  $A_{l,j}$  and  $B_{l,j}$  subsets of elements of matrices  $A$  and  $B$  that initially reside in the local memory of processor  $j$ . The optimal data layout therefore requires that  $A_{l,j}, B_{l,j} \subset Dom(\mathcal{D}_j)$ . However, the schedule does not specify exactly which elements of  $Dom(\mathcal{D}_j)$  should be in  $A_{l,j}$  and  $B_{l,j}$ . As a consequence of the communication pattern § 7.2, each element of  $A_{l,j}$  and  $B_{l,j}$  is communicated to  $g_m, g_n$  processors, respectively. To prevent data reshuffling, we therefore split each of  $Dom(\mathcal{D}_j)$  into  $g_m$  and  $g_n$  smaller blocks, enforcing that consecutive blocks are assigned to processors that communicate first. This is unlike the distributed CARMA implementation [22], which uses the cyclic distribution among processors in the recursion base case and requires local data reshuffling after each communication round. Another advantage of our blocked data layout is a full compatibility with the block-cyclic one, which is used in other linear-algebra libraries.

## 8 EVALUATION

We evaluate COSMA’s communication volume and performance against other state-of-the-art implementations with various combinations of matrix dimensions and memory requirements. These scenarios include both synthetic square matrices, in which all algorithms achieve their peak performance, as well as “flat” (two large dimensions) and real-world “tall-and-skinny” (one large dimension) cases with uneven number of processors.

### Comparison Targets

As a comparison, we use the widely used ScaLAPACK library as provided by Intel MKL (version: 18.0.2.199)<sup>6</sup>, as well as Cyclops Tensor Framework<sup>7</sup>, and the original CARMA implementation<sup>8</sup>. We manually tune ScaLAPACK parameters to achieve its maximum performance. Our experiments showed that on Piz Daint it achieves the highest performance when run with 4 MPI ranks per compute node, 9 cores per rank. Therefore, for each matrix size/node count configuration, we recompute the optimal rank decomposition for ScaLAPACK. Remaining implementations use default decomposition strategy and perform best utilizing 36 ranks per node, 1 core per rank.

### Infrastructure and Implementation Details

All implementations were compiled using the GCC 6.2.0 compiler. We use Cray-MPICH 3.1 implementation of MPI. The parallelism within a rank of ScaLAPACK<sup>9</sup> is handled internally by the MKL BLAS (with GNU OpenMP threading) version 2017.4.196. To profile MPI communication volume, we use the mpiP version 3.4.1 [56].

<sup>6</sup>the latest version available on Piz Daint when benchmarks were performed (August 2018). No improvements of P[S,D,C,Z]GEMM have been reported in the MKL release notes since then.

<sup>7</sup><https://github.com/cyclops-community/ctf>, commit ID 244561c on May 15, 2018

<sup>8</sup><https://github.com/lipshitz/CAPS>, commit ID 7589212 on July 19, 2013

<sup>9</sup>only ScaLAPACK uses multiple cores per ranks



### Experimental Setup and Architectures

We run our experiments on the CPU partition of CSCS Piz Daint, which has 1,813 XC40 nodes with dual-socket Intel Xeon E5-2695 v4 processors ( $2 \cdot 18$  cores, 3.30 GHz, 45 MiB L3 shared cache, 64 GiB DDR3 RAM), interconnected by the Cray Aries network with a dragonfly network topology. We set  $p$  to a number of available cores<sup>10</sup> and  $S$  to the main memory size per core (§ 2.1). To additionally capture cache size per core, the model can be extended to a three-level memory hierarchy. However, cache-size tiling is already handled internally by the MKL.

### Matrix Dimensions and Number of Cores

We use square ( $m = n = k$ ), “largeK” ( $m = n \ll k$ ), “largeM” ( $m \gg n = k$ ), and “flat” ( $m = n \gg k$ ) matrices. The matrix dimensions and number of cores are (1) powers of two  $m = 2^{r_1}$ ,  $n = 2^{r_2}$ ,  $k = 2^{r_3}$ , (2) determined by the real-life simulations or hardware architecture (available nodes on a computer), (3) chosen adversarially, e.g.,  $n^3 + 1$ . Tall matrix dimensions are taken from an application benchmark, namely the calculation of the random phase approximation (RPA) energy of water molecules [21]. There, to simulate  $w$  molecules, the sizes of the matrices are  $m = n = 136w$  and  $k = 228w^2$ . In the strong scaling scenario, we use  $w = 128$  as in the original paper, yielding  $m = n = 17,408$ ,  $k = 3,735,552$ . For performance runs, we scale up to 512 nodes (18,432 cores).

### Selection of Benchmarks

We perform both strong scaling and *memory scaling* experiments. The memory scaling scenario fixes the input size per core ( $\frac{pS}{T}$ ,  $I = mn + mk + nk$ ), as opposed to the work per core ( $\frac{mnk}{p} \neq \text{const}$ ). We evaluate two cases: (1) “limited memory” ( $\frac{pS}{T} = \text{const}$ ), and (2) “extra memory” ( $\frac{p^{2/3}S}{T} = \text{const}$ ).

To provide more information about the impact of communication optimizations on the total runtime, for each of the matrix shapes we also separately measure time spent by COSMA on different parts of the code. For each matrix shape we present two extreme cases of strong scaling - with smallest number of processors (most compute-intensive) and with the largest (most communication-intensive). To additionally increase information provided, we perform these measurements with and without communication-computation overlap.

### Programming Models

We use either the RMA or the Message Passing models. CTF also uses both models, whereas CARMA and ScaLAPACK use MPI two-sided (Message Passing).

### Experimentation Methodology

For each combination of parameters, we perform 5 runs, each with different node allocation. As all the algorithms use BLAS routines for local matrix computations, for each run we execute the kernels three times and take the minimum to compensate for the BLAS setup overhead. We report median and 95% confidence intervals of the runtimes.

## 9 RESULTS

We now present the experimental results comparing COSMA with the existing algorithms. For both strong and memory scaling, we measure total communication volume and runtime on both square

and tall matrices. Our experiments show that COSMA always communicates least data and is the fastest in *all* scenarios.

### Summary and Overall Speedups

As discussed in § 8, we evaluate three benchmarks – strong scaling, “limited memory” (no redundant copies of the input are possible), and “extra memory” ( $p^{1/3}$  extra copies of the input can fit into combined memory of all cores). Each of them we test for square, “largeK”, “largeM”, and “flat” matrices, giving twelve cases in total. In Table 3, we present arithmetic mean of total communication volume per MPI rank across all core counts. We also report the summary of minimum, geometric mean, and maximum speedups vs the second best-performing algorithm.

### Communication Volume

As analyzed in § 5 and § 6, COSMA reaches I/O lower bound (up to the factor of  $\sqrt{S}/(\sqrt{S} + 1)$ ). Moreover, optimizations presented in § 7 secure further improvements compared to other state-of-the-art algorithms. In all cases, COSMA performs least communication. Total communication volume for square and “largeK” scenarios is shown in Figures 6 and 9.

### Square Matrices

Figure 8 presents the % of achieved peak hardware performance for square matrices in all three scenarios. As COSMA is based on the near optimal schedule, it achieves the highest performance *in all cases*. Moreover, its performance pattern is the most stable: when the number of cores is not a power of two, the performance does not vary much compared to all remaining three implementations. We note that matrix dimensions in the strong scaling scenarios ( $m = n = k = 2^{14}$ ) are very small for distributed setting. Yet even in this case COSMA maintains relatively high performance for large numbers of cores: using 4k cores it achieves 35% of peak performance, compared to <5% of CTF and ScaLAPACK, showing excellent strong scaling characteristics.

### Tall and Skinny Matrices

Figure 9 presents the results for “largeK” matrices - due to space constraints, the symmetric “largeM” case is. For strong scaling, the minimum number of cores is 2048 (otherwise, the matrices of size  $m = n = 17,408$ ,  $k = 3,735,552$  do not fit into memory). Again, COSMA shows the most stable performance with a varying number of cores.

### “Flat” Matrices

Matrix dimensions for strong scaling are set to  $m = n = 2^{17} = 131,072$  and  $k = 2^9 = 512$ . Our weak scaling scenario models the rank- $k$  update kernel, with fixed  $k = 256$ , and  $m = n$  scaling accordingly for the “limited” and “extra” memory cases. Such kernels take most of the execution time in, e.g., matrix factorization algorithms, where updating Schur complements is performed as a rank- $k$  gemm operation [31].

### Unfavorable Number of Processors

Due to the processor grid optimization (§ 7.1), the performance is stable and does not suffer from unfavorable combinations of parameters. E.g., the runtime of COSMA for square matrices  $m = n = k = 16,384$  on  $p_1 = 9,216 = 2^{10} \cdot 3^2$  cores is 142 ms. Adding an extra core ( $p_2 = 9,217 = 13 \cdot 709$ ), does not change COSMA’s runtime, as the optimal decomposition does not utilize it. On the other hand, CTF for  $p_1$  runs in 600 ms, while for  $p_2$  the runtime *increases* to 1613 ms due to a non-optimal processor decomposition.

<sup>10</sup>for ScaLAPACK, actual number of MPI ranks is  $p/9$

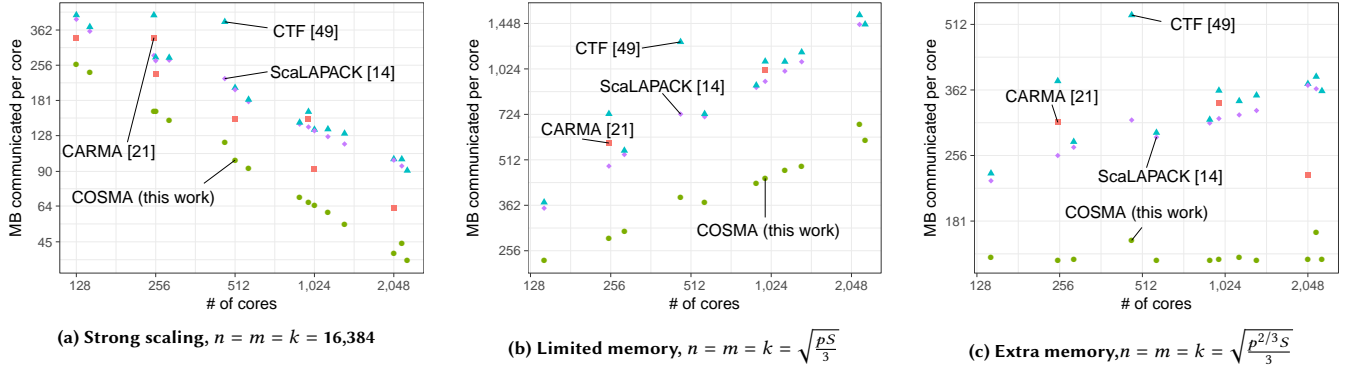


Figure 6: Total communication volume per core carried out by COSMA, CTF, ScaLAPACK and CARMA for square matrices, as measured by the mpiP profiler.

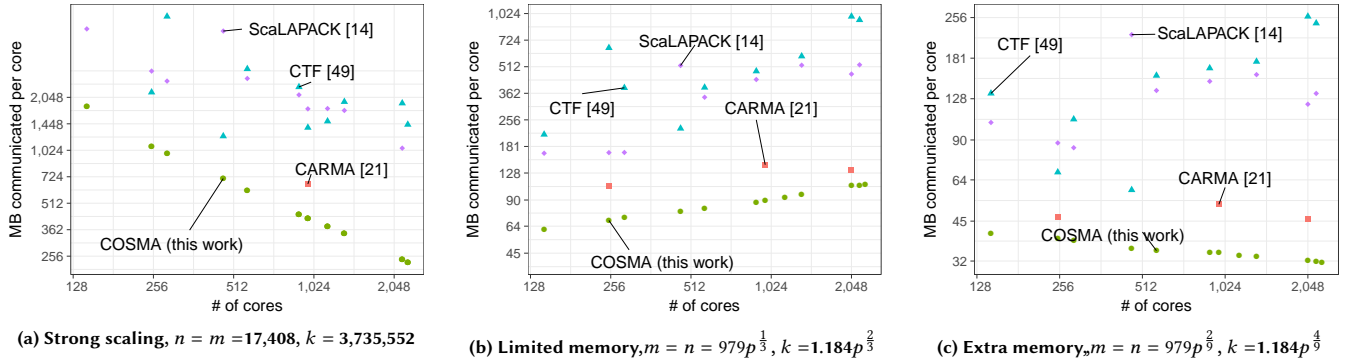


Figure 7: Total communication volume per core carried out by COSMA, CTF, ScaLAPACK and CARMA for “largeK” matrices, as measured by the mpiP profiler.

### Communication-Computation Breakdown

In Figure 10 we present the total runtime breakdown of COSMA into communication and computation routines. Combined with the comparison of communication volumes (Figures 6 and 7, Table 3) we see the importance of our I/O optimizations for distributed setting even for traditionally compute-bound MMM. E.g., for square or “flat” matrix and 16k cores, COSMA communicates more than two times less than the second-best (CARMA). Assuming constant time-per-MB, COSMA would be 40% slower if it communicated that much, being slower than CARMA by 30%. For “largeK”, the situation is even more extreme, with COSMA suffering 2.3 times slowdown if communicating as much as the second-best algorithm, CTF, which communicates 10 times more.

### Detailed Statistical Analysis

Figure 11 provides a distribution of the achieved peak performance across all numbers of cores for all six scenarios. It can be seen that, for example, in the strong scaling scenario and square matrices, COSMA is comparable to the other implementations (especially CARMA). However, for tall-and-skinny matrices with limited memory available, COSMA lowest achieved performance is higher than the best performance of CTF and ScaLAPACK.

## 10 RELATED WORK

Works on data movement minimization may be divided into two categories: applicable across memory hierarchy (vertical, also called I/O minimization), or between parallel processors (horizontal, also called communication minimization). Even though they are “two


shape	benchmark	total comm. volume per rank [MB]				speedup		
		ScaLAPACK	CTF	CARMA	COSMA	min	mean	max
B A C	strong scaling	203	222	195	107	1.07	1.94	4.81
	limited memory	816	986	799	424	1.23	1.71	2.99
	extra memory	303	350	291	151	1.14	2.03	4.73
B A C	strong scaling	2636	2278	659	545	1.24	2.00	6.55
	limited memory	368	541	128	88	1.30	2.61	8.26
	extra memory	133	152	48	35	1.31	2.55	6.70
B A C	strong scaling	3507	2024	541	410	1.31	2.22	3.22
	limited memory	989	672	399	194	1.42	1.7	2.27
	extra memory	122	77	77	29	1.35	1.76	2.8
B A C	strong scaling	134	68	10	7	1.21	4.02	12.81
	limited memory	47	101	26	8	1.31	2.07	3.41
	extra memory	15	15	10	3	1.5	2.29	3.59
overall						1.07	2.17	12.81

Table 3: Average communication volume per MPI rank and measured speedup of COSMA vs the second-best algorithm across all core counts for each of the scenarios.


sides of the same coin”, in literature they are often treated as separate topics. In our work we combine them: analyze trade-offs between communication optimal (distributed memory) and I/O optimal schedule (shared memory).

### 10.1 General I/O Lower Bounds


Hong and Kung [34] analyzed the I/O complexity for general CDAGs in their the red-blue pebble game, on which we base our work. As a special case, they derived an asymptotic bound  $\Omega\left(n^3/\sqrt{S}\right)$  for MMM. Elango et al. [23] extended this work to the red-blue-white game and Liu and Terman [40] proved that it is also P-SPACE complete. Irony et al. [33] extended the MMM lower bound result to



**Figure 8:** Achieved % of peak performance by COSMA, CTF, ScaLAPACK and CARMA for square matrices, strong and weak scaling. We show median and 95% confidence intervals.



**Figure 9:** Achieved % of peak performance by COSMA, CTF, ScaLAPACK and CARMA for “largeK” matrices, strong and weak scaling. We show median and 95% confidence intervals.



**Figure 10:** Time distribution of COSMA communication and computation kernels for strong scaling executed on the smallest and the largest core counts for each of the matrix shapes. Left bar: no communication–computation overlap. Right bar: overlap enabled.

a parallel machine with  $p$  processors, each having a fast private memory of size  $S$ , proving the  $\frac{n^3}{2\sqrt{2}p\sqrt{S}} - S$  lower bound on the communication volume per processor. Chan [12] studied different variants of pebble games in the context of memory space and parallel time. Aggarwal and Vitter [2] introduced a two-memory machine that models a blocked access and latency in an external storage. Arge et al. [3] extended this model to a parallel machine. Solomonik et al. [50] combined the communication, synchronization, and computation in their general cost model and applied it to several linear algebra algorithms. Smith and van de Geijn [47] derived a sequential lower bound  $2mnk/\sqrt{S} - 2S$  for MMM. They showed that the leading factor  $2mnk/\sqrt{S}$  is tight. We improve this result by 1) improving an additive factor of  $2S$ , but more importantly

2) generalizing the bound to a parallel machine. Our work uses a simplified model, not taking into account the memory block size, as in the external memory model, nor the cost of computation. We motivate it by assuming that the block size is significantly smaller than the input size, the data is layout contiguously in the memory, and that the computation is evenly distributed among processors.

## 10.2 Shared Memory Optimizations

I/O optimization for linear algebra includes such techniques as loop tiling and skewing [58], interchanging and reversal [57]. For programs with multiple loop nests, Kennedy and McKinley [35] showed various techniques for loop fusion and proved that in general this problem is NP-hard. Later, Darté [20] identified cases when this problem has polynomial complexity.

Toledo [54] in his survey on Out-Of-Core (OOC) algorithms analyzed various I/O minimizing techniques for dense and sparse matrices. Mohanty [42] in his thesis optimized several OOC algorithms. Irony et al. [33] proved the I/O lower bound of classical MMM on a parallel machine. Ballard et al. [5] proved analogous results for Strassen’s algorithm. This analysis was extended by Scott et al. [45] to a general class of Strassen-like algorithms.

Although we consider only dense matrices, there is an extensive literature on sparse matrix I/O optimizations. Bender et al. [7] extended Aggarwal’s external memory model [2] and showed I/O complexity of the sparse matrix-vector (SpMV) multiplication. Greiner [29] extended those results and provided I/O complexities of other sparse computations.



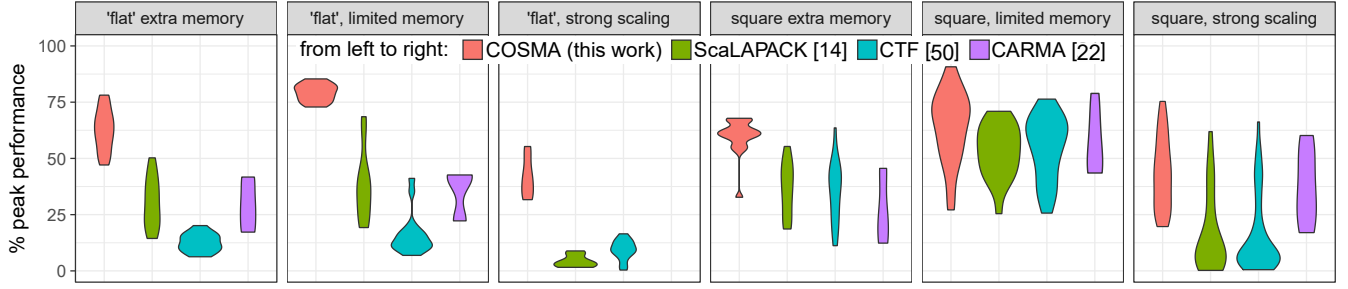


Figure 11: Distribution of achieved % of peak performance of the algorithms across all number of cores for “flat” and square matrices.

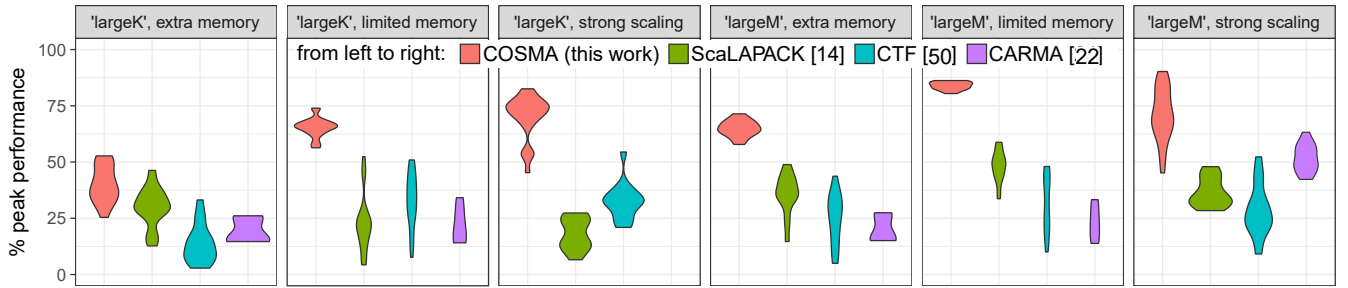


Figure 12: Distribution of achieved % of peak performance of the algorithms across all number of cores for tall-and-skinny matrices.

### 10.3 Distributed Memory Optimizations

Distributed algorithms for dense matrix multiplication date back to the work of Cannon [10], which has been analyzed and extended many times [30] [39]. In the presence of extra memory, Aggarwal et al. [1] included parallelization in the third dimension. Solomonik and Demmel [52] extended this scheme with their 2.5D decomposition to arbitrary range of the available memory, effectively interpolating between Cannon’s 2D and Aggarwal’s 3D scheme. A recursive, memory-oblivious MMM algorithm was introduced by Blumofe et al. [9] and extended to rectangular matrices by Frigo et al. [26]. Demmel et al. [22] introduced CARMA algorithm which achieves the asymptotic complexity for all matrix and memory sizes. We compare COSMA with these algorithms, showing that we achieve better results both in terms of communication complexity and the actual runtime performance. Lazzaro et al. [38] used the 2.5D technique for sparse matrices, both for square and rectangular grids. Koanantakool et al. [37] observed that for sparse-dense MMM, 1.5D decomposition performs less communication than 2D and 2.5D schemes, as it distributes only the sparse matrix.

## 11 CONCLUSIONS

In this work we present a new method (Lemma 1) for assessing tight I/O lower bounds of algorithms using their CDAG representation and the red-blue pebble game abstraction. As a use case, we prove a tight bound for MMM, both for a sequential (Theorem 1) and parallel (Theorem 2) execution. Furthermore, our proofs are constructive: our COSMA algorithm is near I/O optimal (up to the factor of  $\frac{\sqrt{S}}{\sqrt{S+1}-1}$ , which is less than 0.04% from the lower bound for 10MB of fast memory) for any combination of matrix dimensions, number of

processors and memory sizes. This is in contrast with the current state-of-the-art algorithms, which are communication-inefficient in some scenarios.

To further increase the performance, we introduce a series of optimizations, both on an algorithmic level (processor grid optimization (§ 7.1) and blocked data layout (§ 7.6)) and hardware-related (enhanced communication pattern (§ 7.2), communication-computation overlap (§ 7.3), one-sided (§ 7.4) communication). The experiments confirm the superiority of COSMA over the other analyzed algorithms - our algorithm significantly reduces communication in all tested scenarios, supporting our theoretical analysis. Most importantly, our work is of practical importance, being maintained as an open-source implementation and achieving a time-to-solution speedup of up to 12.8x times compared to highly optimized state-of-the-art libraries.

The important feature of our method is that it does not require any manual parameter tuning and is generalizable to other machine models (e.g., multiple levels of memory) and linear algebra kernels (e.g., LU or Cholesky decompositions), both for dense and sparse matrices. We believe that the “bottom-up” approach will lead to developing more efficient distributed algorithms in the future.

### Acknowledgements

We thank Yishai Oltchik and Niels Gleinig for invaluable help with the theoretical part of the paper, and Simon Pintarelli for advice and support with the implementation. We also thank CSCS for the compute hours needed to conduct all experiments. This project has received funding from the European Research Council (ERC) under the European Union’s Horizon2020 programme (grant agreement DAPP, No.678880), and additional funding from the Platform for Advanced Scientific Computing (PASC).

## REFERENCES

- [1] R. C. Agarwal et al. 1995. A three-dimensional approach to parallel matrix multiplication. *IBM J. Res. Dev.* (1995).
- [2] Alok Aggarwal and S. Vitter, Jeffrey. [n. d.]. The Input/Output Complexity of Sorting and Related Problems. *CACM* (Sept. [n. d.]).
- [3] Lars Arge et al. 2008. In *SPAA*.
- [4] Ariful Azad et al. 2015. Parallel triangle counting and enumeration using matrix algebra. In *IPDPS*.
- [5] Grey Ballard et al. 2012. Graph expansion and communication costs of fast matrix multiplication. *JACM* (2012).
- [6] Tal Ben-Nun and Torsten Hoefler. 2018. Demystifying Parallel and Distributed Deep Learning: An In-Depth Concurrency Analysis. *CoRR* abs/1802.09941 (2018).
- [7] Michael A Bender et al. 2010. Optimal sparse matrix dense vector multiplication in the I/O-model. *TOCS* (2010).
- [8] Maciej Besta et al. 2017. SlimSell: A Vectorizable Graph Representation for Breadth-First Search. In *IPDPS*.
- [9] Robert D Blumofe et al. 1996. An analysis of dag-consistent distributed shared-memory algorithms. In *SPAA*.
- [10] Lynn Elliot Cannon. 1969. *A Cellular Computer to Implement the Kalman Filter Algorithm*. Ph.D. Dissertation.
- [11] Gregory J. Chaitin et al. 1981. Register allocation via coloring. *Computer Languages* (1981).
- [12] S. M. Chan. 2013. Just a Pebble Game. In *CCC*.
- [13] Françoise Chatelin. 2012. *Eigenvalues of Matrices: Revised Edition*. Siam.
- [14] Jaeyoung Choi et al. 1992. ScaLAPACK: A scalable linear algebra library for distributed memory concurrent computers. In *FRONTIERS*.
- [15] Jaeyoung Choi et al. 1996. Design and Implementation of the ScaLAPACK LU, QR, and Cholesky Factorization Routines. *Sci. Program.* (1996).
- [16] J. Choi et al. 1996. ScaLAPACK: a portable linear algebra library for distributed memory computers — design issues and performance. *Comp. Phys. Comm.* (1996).
- [17] Jaeyoung Choi, David W Walker, and Jack J Dongarra. 1994. PUMMA: Parallel universal matrix multiplication algorithms on distributed memory concurrent computers. *Concurrency: Practice and Experience* 6, 7 (1994), 543–570.
- [18] Thomas H Cormen, Charles E Leiserson, Ronald L Rivest, and Clifford Stein. 2009. *Introduction to algorithms*. MIT press.
- [19] Paolo D'Alberto and Alexandru Nicolau. 2008. Using recursion to boost ATLAS's performance. In *High-Performance Computing*. Springer.
- [20] Alain Darté. 1999. On the complexity of loop fusion. In *PACT*.
- [21] Mauro Del Ben et al. 2015. Enabling simulation at the fifth rung of DFT: Large scale RPA calculations with excellent time to solution. *Comp. Phys. Comm.* (2015).
- [22] J. Demmel et al. 2013. Communication-Optimal Parallel Recursive Rectangular Matrix Multiplication. In *IPDPS*.
- [23] V. Elango et al. 2013. *Data access complexity: The red/blue pebble game revisited*. Technical Report.
- [24] Geoffrey C Fox. 1988. Solving problems on concurrent processors. (1988).
- [25] Geoffrey C Fox, Steve W Otto, and Anthony JG Hey. 1987. Matrix algorithms on a hypercube I: Matrix multiplication. *Parallel computing* 4, 1 (1987), 17–31.
- [26] M. Frigo et al. 1999. Cache-oblivious algorithms. In *FOCS*.
- [27] R. Gerstenberger et al. 2013. Enabling Highly-Scalable Remote Memory Access Programming with MPI-3 One Sided. In *SC*.
- [28] John R. Gilbert et al. 1979. The Pebbling Problem is Complete in Polynomial Space. In *STOC*.
- [29] Gero Greiner. 2012. *Sparse matrix computations and their I/O complexity*. Ph.D. Dissertation. Technische Universität München.
- [30] A. Gupta and V. Kumar. 1993. Scalability of Parallel Algorithms for Matrix Multiplication. In *ICPP*.
- [31] Azzam Haidar, Stanimire Tomov, Jack Dongarra, and Nicholas J Higham. 2018. Harnessing GPU tensor cores for fast FP16 arithmetic to speed up mixed-precision iterative refinement solvers. In *Proceedings of the International Conference for High Performance Computing, Networking, Storage, and Analysis*. IEEE Press, 47.
- [32] T. Hoefler et al. 2015. Remote Memory Access Programming in MPI-3. *TOPC* (2015).
- [33] Dror Irony et al. 2004. Communication Lower Bounds for Distributed-memory Matrix Multiplication. *JPDC* (2004).
- [34] Hong Jia-Wei and Hsiang-Tsung Kung. 1981. I/O complexity: The red-blue pebble game. In *STOC*.
- [35] Ken Kennedy and Kathryn S McKinley. 1993. Maximizing loop parallelism and improving data locality via loop fusion and distribution. In *LCPC*.
- [36] Jeremy Kepner et al. 2016. Mathematical foundations of the GraphBLAS. *arXiv:1606.05790* (2016).
- [37] Penporn Koanantakool et al. 2016. Communication-avoiding parallel sparse-dense matrix-matrix multiplication. In *IPDPS*.
- [38] Alfio Lazzaro et al. 2017. Increasing the efficiency of sparse matrix-matrix multiplication with a 2.5 D algorithm and one-sided MPI. In *PASC*.
- [39] Hyuk-Jae Lee et al. 1997. Generalized Cannon's Algorithm for Parallel Matrix Multiplication. In *ICS*.
- [40] Quanquan Liu. 2018. Red-Blue and Standard Pebble Games : Complexity and Applications in the Sequential and Parallel Models.
- [41] Carl D Meyer. 2000. *Matrix analysis and applied linear algebra*. SIAM.
- [42] Sraban Kumar Mohanty. 2010. I/O Efficient Algorithms for Matrix Computations. *CoRR* abs/1006.1307 (2010).
- [43] Andrew Y Ng et al. 2002. On spectral clustering: Analysis and an algorithm. In *NIPS*.
- [44] Donald W. Rogers. 2003. *Computational Chemistry Using the PC*. John Wiley & Sons, Inc.
- [45] Jacob Scott et al. 2015. Matrix multiplication I/O-complexity by path routing. In *SPAA*.
- [46] Ravi Sethi. 1973. Complete Register Allocation Problems. In *STOC*.
- [47] Tyler Michael Smith and Robert A. van de Geijn. 2017. Pushing the Bounds for Matrix-Matrix Multiplication. *CoRR* abs/1702.02017 (2017).
- [48] Raffaele Solcà, Anton Kozhevnikov, Azzam Haidar, Stanimire Tomov, Jack Dongarra, and Thomas C. Schulthess. 2015. Efficient Implementation of Quantum Materials Simulations on Distributed CPU-GPU Systems. In *Proceedings of the International Conference for High Performance Computing, Networking, Storage and Analysis (SC '15)*. ACM, New York, NY, USA, Article 10, 12 pages. <https://doi.org/10.1145/2807591.2807654>
- [49] Edgar Solomonik et al. 2013. Cyclops tensor framework: Reducing communication and eliminating load imbalance in massively parallel contractions. In *IPDPS*.
- [50] Edgar Solomonik et al. 2016. Trade-Offs Between Synchronization, Communication, and Computation in Parallel Linear Algebra computations. *TOPC* (2016).
- [51] E. Solomonik et al. 2017. Scaling Betweenness Centrality using Communication-Efficient Sparse Matrix Multiplication. In *SC*.
- [52] Edgar Solomonik and James Demmel. 2011. Communication-Optimal Parallel 2.5D Matrix Multiplication and LU Factorization Algorithms. In *EuroPar*.
- [53] Volker Strassen. 1969. Gaussian Elimination is Not Optimal. *Numer. Math.* (1969).
- [54] Sivan Toledo. 1999. A survey of out-of-core algorithms in numerical linear algebra. *External Memory Algorithms and Visualization* (1999).
- [55] Robert A Van De Geijn and Jerrell Watts. 1997. SUMMA: Scalable universal matrix multiplication algorithm. *Concurrency: Practice and Experience* 9, 4 (1997), 255–274.
- [56] Jeffrey S Vetter and Michael O McCracken. 2001. Statistical scalability analysis of communication operations in distributed applications. *ACM SIGPLAN Notices* 36, 7 (2001), 123–132.
- [57] Michael E. Wolf and Monica S. Lam. 1991. A Data Locality Optimizing Algorithm. In *PLDI*.
- [58] M. Wolfe. 1989. More Iteration Space Tiling. In *SC*.
- [59] Nan Xiong. 2018. *Optimizing Tall-and-Skinny Matrix-Matrix Multiplication on GPUs*. Ph.D. Dissertation. UC Riverside.
- [60] Qinqing Zheng and John D. Lafferty. 2016. Convergence Analysis for Rectangular Matrix Completion Using Burer-Monteiro Factorization and Gradient Descent. *CoRR* (2016).

# Appendix: Artifact Description/Artifact Evaluation

## SUMMARY OF THE EXPERIMENTS REPORTED

We evaluated COSMA against: -ScaLAPACK as provided by Intel MKL (version: 18.0.2.199), -Cyclops Tensor Framework (<https://github.com/cyclops-community/ctf>, commit ID 244561c on May 15, 2018), -the original CARMA implementation (<https://github.com/lipshitz/CAPS> commit ID 7589212 on July 19, 2013).

We measured both total runtime and communication volume using mpiP profiler version 3.4.1 for 12 matrix multiplication scenarios: -"largeK" matrices, strong scaling -"largeK" matrices, memory scaling p0 -"largeK" matrices, memory scaling p1 -"largeM" matrices, strong scaling -"largeM" matrices, memory scaling p0 -"largeM" matrices, memory scaling p1 -"flat" matrices, strong scaling -"flat" matrices, memory scaling p0 -"flat" matrices, memory scaling p1 -square matrices, strong scaling -square matrices, memory scaling p0 -square matrices, memory scaling p1,

All our experiments were run on the CPU partition of the CSCS Piz Daint supercomputer, equipped with 1,813 XC40 nodes with dual-socket Intel Xeon E5-2695 v4 processors (3.30 GHz, 45 MiB L3 shared cache, 64 GiB DDR3 RAM), interconnected with Cray Aries network.

All implementations were compiled using the gcc 6.2.0 compiler. We use Cray-MPICH 3.1 implementation of MPI. The intra-node parallelism is handled internally by the MKL BLAS version 2017.4.196.

## ARTIFACT AVAILABILITY

*Software Artifact Availability:* All author-created software artifacts are maintained in a public repository under an OSI-approved license.

*Hardware Artifact Availability:* There are no author-created hardware artifacts.

*Data Artifact Availability:* There are no author-created data artifacts.

*Proprietary Artifacts:* None of the associated artifacts, author-created or otherwise, are proprietary.

*List of URLs and/or DOIs where artifacts are available:*

<https://github.com/eth-cscs/COSMA>  
[https://www.dropbox.com/sh/u7b9fk1uvb6t04t/AAAMfo4yf\\_40iwOuLer\\_Zvvz5a?dl=0](https://www.dropbox.com/sh/u7b9fk1uvb6t04t/AAAMfo4yf_40iwOuLer_Zvvz5a?dl=0) technical  
↪ review

## BASELINE EXPERIMENTAL SETUP, AND MODIFICATIONS MADE FOR THE PAPER

*Relevant hardware details:* CPU partition of the CSCS Piz Daint supercomputer, equipped with 1,813 XC40 nodes with dual-socket Intel Xeon E5-2695 v4

*Operating systems and versions:* SUSE Linux Enterprise Server 12 SP3

*Compilers and versions:* gcc 6.2.0

*Applications and versions:* ScaLAPACK as provided by Intel MKL (version: 17.0.4.196), Cyclops Tensor Framework, CARMA

*Libraries and versions:* Cray-MPICH 3.1, MKL BLAS version 2017.4.196

*Key algorithms:* matrix multiplication

*Input datasets and versions:* synthetic matrices

*Paper Modifications:* We manually tune ScaLAPACK parameters. Our experiments showed that on Piz Daint it achieves the highest performance when run with 4 ranks per compute node, 9 threads per rank. Therefore, for each matrix sizes/processor count configuration, we recompute optimal processor grid for ScaLAPACK.

*Output from scripts that gathers execution environment information.*

```
LESSKEY=/etc/lesskey.bin
MODULE_VERSION_STACK=3.2.10.6
KSH_AUTOLOAD=1
CRAY_BINUTILS_BIN=/opt/cray/pe/cce/8.7.3/binutils/x8_
↪ 6_64/bin
PE_LIBSCI_VOLATILE_PRGENV=CRAY GNU INTEL
PE_SMA_DEFAULT_PKGCONFIG_VARIABLES=PE_SMA_COMPFLAG_@
↪ prgenv@
PE_TPSSL_64_DEFAULT_GENCOMPS_INTEL_mic_knl=160
NNTPSERVER=news
PE_PAPI_DEFAULT_ACCEL_FAMILY_LIBS_nvidia=-lcupti,-l_
↪ cudart,-lcuda
PE_PETSC_DEFAULT_GENCOMPILERS_CRAY_sandybridge=8.6
PE_PETSC_DEFAULT_GENCOMPS_CRAY_skylake=86
PE_TPSSL_DEFAULT_GENCOMPS_INTEL_x86_skylake=160
PE_CXX_PKGCONFIG_LIBS=mpichcxx
PE_MPICH_GENCOMPILERS_PGI=15.3
XDG_SESSION_ID=242719
XALT_ETC_DIR=/apps/daint/UES/xalt/0.7.6/etc
HOSTNAME=nid00008
CRAY_UDREG_INCLUDE_OPTS=-I/opt/cray/udreg/2.3.2-6.0_
↪ 7.0_33.18_g5196236.ari/include
GCC_AARCH64=/opt/gcc-cross-aarch64/6.1.0/aarch64
GCC_X86_64=/opt/gcc/6.1.0/snos
PE_FFTW_DEFAULT_TARGET_mic_knl=mic_knl
PE_LIBSCI_ACC_DEFAULT_PKGCONFIG_VARIABLES=PE_LIBSCI_
↪ ACC_DEFAULT_NV_SUFFIX_@accelerator@
PE_PETSC_DEFAULT_GENCOMPILERS_INTEL_mic_knl=16.0
PE_TPSSL_64_DEFAULT_GENCOMPS_INTEL_interlagos=160
PE_TRILINOS_DEFAULT_GENCOMPS_CRAY_x86_64=86
RCLOCAL_BASEOPTS=true
XKEYSYMDB=/usr/X11R6/lib/X11/XKeysymDB
CRAY_SITE_LIST_DIR=/etc/opt/cray/pe/modules
```



```

LIBRARYMODULES=acml:alps:cray-dwarf:cray-fftw:cray-g
↳ a:cray-hdf5:cray-hdf5-parallel:cray-libsci:cray-
↳ libsci_acc:cray-mpich:cray-mpich2:cray-mpich-abi
↳ :cray-netcdf:cray-netcdf-hdf5parallel:cray-paral
↳ lel-netcdf:cray-petsc:cray-petsc-complex:cray-sh
↳ mem:cray-tpsl:cray-trilinos:cuda-toolkit:fftw:ga:
↳ hdf5:hdf5-parallel:iobuf:libfast:netcdf:netcdf-h
↳ df5parallel:ntk:onesided:papi:petsc:petsc-comple
↳ x:pmi:tpsl:trilinos:xt-libsci:xt-mpich2:xt-mpt:x
↳ t-papi
ASSEMBLER_AARCH64=/opt/cray/pe/cce/8.7.3/binutils/cr
↳ oss/x86_64-aarch64/aarch64-linux-gnu/bin/as
PE_NETCDF_DEFAULT_VOLATILE_PKGCONFIG_PATH=/opt/cray/
↳ pe/netcdf/4.6.1.2/@PRGENV@/@PE_NETCDF_DEFAULT_GE
↳ NCOMPS@/lib/pkgconfig
PE_PARALLEL_NETCDF_DEFAULT_VOLATILE_PKGCONFIG_PATH=/
↳ opt/cray/pe/parallel-netcdf/1.8.1.3/@PRGENV@/@PE
↳ _PARALLEL_NETCDF_DEFAULT_GENCOMPS@/lib/pkgconfig
PE_SMA_DEFAULT_COMPFLAG_GNU=-fcray-pointer
PE_TRILINOS_DEFAULT_VOLATILE_PKGCONFIG_PATH=/opt/cra
↳ y/pe/trilinos/12.12.1.0/@PRGENV@/@PE_TRILINOS_DE
↳ FAULT_GENCOMPS@/@PE_TRILINOS_DEFAULT_TARGET@/lib
↳ /pkgconfig
CRAY_BINUTILS_ROOT=/opt/cray/pe/cce/8.7.3/binutils/x
↳ 86_64/x86_64-pc-linux-gnu/./
CRAY_FTN_VERSION=8.7.3
PE_ENV=CRAY
PE_HDF5_DEFAULT_GENCOMPILERS_GNU=7.1 6.1 5.3 4.9
PE_MPICH_ALTERNATE_LIBS_dpm=_dpm
PE_SMA_DEFAULT_COMPFLAG=
PE_TPSL_64_DEFAULT_GENCOMPILERS_CRAY_x86_64=8.6
SHELL=/usr/local/bin/bash
TERM=xterm-256color
HOST=daint106
ASSEMBLER_X86_64=/opt/cray/pe/cce/8.7.3/binutils/x86
↳ _64/x86_64-pc-linux-gnu/bin/as
PE_TPSL_DEFAULT_GENCOMPS_CRAY_x86_skylake=86
PKGCONFIG_ENABLED=1
HISTSIZE=1000
PROFILEREAD=true
LINKER_AARCH64=/opt/cray/pe/cce/8.7.3/binutils/cross
↳ /x86_64-aarch64/aarch64-linux-gnu/bin/ld
PE_PETSC_DEFAULT_GENCOMPS_CRAY_sandybridge=86
PE_TPSL_DEFAULT_GENCOMPILERS_GNU_x86_skylake=7.1 6.1
TMPDIR=/tmp
CRAYPE_DIR=/opt/cray/pe/craype/2.5.15
CRAY_UGNI_POST_LINK_OPTS=-L/opt/cray/ugni/6.0.14.0-6
↳ .0.7.0-23.1__gea11d3d.ari/lib64
CRAY_XPMEM_POST_LINK_OPTS=-L/opt/cray/xpmem/2.2.15-6
↳ .0.7.1.5.10__g7549d06.ari/lib64
FORTRAN_SYSTEM_MODULE_NAMES=ftn_lib_definitions
PE_NETCDF_DEFAULT_VOLATILE_PRGENV=GNU
PE_PARALLEL_NETCDF_DEFAULT_VOLATILE_PRGENV=GNU
PE_PETSC_DEFAULT_GENCOMPS_GNU_haswell=71 53 49
PE_PETSC_DEFAULT_GENCOMPS_INTEL_haswell=160
PE_TPSL_64_DEFAULT_GENCOMPS_INTEL_x86_skylake=160

```

```

PE_TPSL_DEFAULT_GENCOMPS_GNU_sandybridge=71 53 49
PE_TPSL_DEFAULT_REQUIRED_PRODUCTS=PE_MPICH:PE_LIBSCI
PE_TRILINOS_DEFAULT_VOLATILE_PRGENV=CRAY GNU INTEL
PE_MPICH_DIR_PGI_DEFAULT64=64
PE_FFTW_DEFAULT_VOLATILE_PKGCONFIG_PATH=/opt/cray/pe
↳ /fftw/3.3.6.5/@PE_FFTW_DEFAULT_TARGET@/lib/pkgco
↳ nfig
PE_HDF5_DEFAULT_VOLATILE_PRGENV=GNU
PE_HDF5_PARALLEL_DEFAULT_VOLATILE_PKGCONFIG_PATH=/op
↳ t/cray/pe/hdf5-parallel/1.10.2.0/@PRGENV@/@PE_HD
↳ F5_PARALLEL_DEFAULT_GENCOMPS@/lib/pkgconfig
PE_NETCDF_HDF5PARALLEL_DEFAULT_VOLATILE_PKGCONFIG_PA
↳ TH=/opt/cray/pe/netcdf-hdf5parallel/4.6.1.2/@PRG
↳ ENV@/@PE_NETCDF_HDF5PARALLEL_DEFAULT_GENCOMPS@/l
↳ ib/pkgconfig
PE_PETSC_DEFAULT_GENCOMPS_CRAY_interlagos=86
CRAY_MPICH2_DIR=/opt/cray/pe/mpt/7.7.2/gni/mpich-cra
↳ y/8.6
ALT_LINKER=/apps/daint/UES/xalt/0.7.6/bin/ld
PMI_CONTROL_PORT=22748
PE_GA_DEFAULT_VOLATILE_PRGENV=GNU
PE_LIBSCI_DEFAULT_GENCOMPS_GNU_x86_64=71 61 51 49
PE_TPSL_64_DEFAULT_GENCOMPILERS_CRAY_interlagos=8.6
PE_TPSL_DEFAULT_GENCOMPS_CRAY_mic_knl=86
MORE=-s1
FPATH=/opt/cray/pe/modules/3.2.10.6/init/sh_funcs/n
↳ o_redirect:/opt/cray/pe/modules/3.2.10.6/init/sh
↳ _funcs/no_redirect
PERFTOOLS_VERSION=7.0.3
PE_LIBSCI_ACC_DEFAULT_GENCOMPS_CRAY_x86_64=85
PE_LIBSCI_ACC_DEFAULT_REQUIRED_PRODUCTS=PE_MPICH:PE_
↳ LIBSCI
PE_MPICH_DEFAULT_GENCOMPILERS_GNU=7.1 5.1 4.9
PE_PKGCONFIG_PRODUCTS=PE_MPICH:PE_LIBSCI
PE_TPSL_DEFAULT_GENCOMPS_INTEL_x86_64=160
PE_MPICH_GENCOMPS_GNU=71 51 49
PE_PAPI_DEFAULT_ACCEL_LIBS_nvidia35=-lcupti,-lcudar
↳ t,-lcuda
PE_PETSC_DEFAULT_REQUIRED_PRODUCTS=PE_MPICH:PE_LIBSC
↳ I:PE_HDF5_PARALLEL:PE_TPSL
PE_TPSL_64_DEFAULT_GENCOMPS_CRAY_haswell=86
PE_TPSL_64_DEFAULT_VOLATILE_PKGCONFIG_PATH=/opt/cray
↳ /pe/tpsl/18.06.1/@PRGENV@64/@PE_TPSL_64_DEFAULT_
↳ GENCOMPS@/@PE_TPSL_64_DEFAULT_TARGET@/lib/pkgcon
↳ fig
PE_CRAY_DEFAULT_FIXED_PKGCONFIG_PATH=/opt/cray/pe/pa
↳ rallel-netcdf/1.8.1.3/CRAY/8.6/lib/pkgconfig:/op
↳ t/cray/pe/netcdf-hdf5parallel/4.6.1.2/CRAY/8.6/l
↳ ib/pkgconfig:/opt/cray/pe/netcdf/4.6.1.2/CRAY/8.
↳ 6/lib/pkgconfig:/opt/cray/pe/hdf5-parallel/1.10.
↳ 2.0/CRAY/8.6/lib/pkgconfig:/opt/cray/pe/hdf5/1.1
↳ 0.2.0/CRAY/8.6/lib/pkgconfig:/opt/cray/pe/ga/5.3
↳ .0.8/CRAY/8.6/lib/pkgconfig
PE_TRILINOS_DEFAULT_GENCOMPILERS_CRAY_x86_64=8.6
PE_LIBSCI_DEFAULT_OMP_REQUIRES_openmp=_mp
PE_PETSC_DEFAULT_GENCOMPS_CRAY_x86_64=86

```

## Red-Blue Pebbling Revisited: Near Optimal Parallel Matrix Multiplication

```
PE_TPSL_64_DEFAULT_GENCOMPILERS_CRAY_sandybridge=8.6
cce_already_loaded=0
PE_FORTRAN_PKGCONFIG_LIBS=mpichf90
CRAYPAT_ALPS_COMPONENT=/opt/cray/pe/perftools/7.0.3/
↳ sbin/pat_alps
CRAYPAT_LD_LIBRARY_PATH=/opt/cray/pe/gcc-libs:/opt/c
↳ ray/gcc-libs:/opt/cray/pe/perftools/7.0.3/lib64
CRAY_BINUTILS_ROOT_AARCH64=/opt/cray/pe/cce/8.7.3/bi
↳ nutils/cross/x86_64-aarch64/aarch64-linux-gnu/./
CRAY_BINUTILS_VERSION=/opt/cray/pe/cce/8.7.3
CRAY_PRGENVCRAY=loaded
PE_SMA_DEFAULT_VOLATILE_PKGCONFIG_PATH=/opt/cray/pe/
↳ mpt/7.7.2/gni/sma@PE_SMA_DEFAULT_DIR_DEFAULT64@/
↳ lib64/pkgconfig
ALLINEA_QUEUE_DLL=/opt/cray/pe/mpt/7.7.2/gni/mpich-c
↳ ray/8.6/lib/libtvmmpich.so.3.0.1
ALPS_APP_ID=12841473
PE_LIBSCI_ACC_DEFAULT_VOLATILE_PRGENV=CRAY GNU
PE_TRILINOS_DEFAULT_GENCOMPS_INTEL_x86_64=160
CRAY_MPICH_BASEDIR=/opt/cray/pe/mpt/7.7.2/gni
JRE_HOME=/usr/lib64/jvm/java/jre
PE_HDF5_PARALLEL_DEFAULT_GENCOMPILERS_GNU=7.1 6.1 5.3
↳ 4.9
PE_NETCDF_HDF5PARALLEL_DEFAULT_GENCOMPILERS_GNU=7.1
↳ 6.1 5.3 4.9
PE_TPSL_64_DEFAULT_GENCOMPS_CRAY_x86_skylake=86
PE_TPSL_64_DEFAULT_GENCOMPS_INTEL_haswell=160
LS_COLORS=no=00:fi=00:di=01;34:ln=00;36:pi=40;33:so=
↳ 01;35:do=01;35:bd=40;33;01:cd=40;33;01:or=41;33;
↳ 01:ex=00;32:*.cmd=00;32:*.exe=01;32:*.com=01;32:
↳ *.bat=01;32:*.btm=01;32:*.dll=01;32:*.tar=00;31:
↳ *.tbz=00;31:*.tgz=00;31:*.rpm=00;31:*.deb=00;31:
↳ *.arj=00;31:*.taz=00;31:*.lzh=00;31:*.lzma=00;31
↳ :*.zip=00;31:*.zoo=00;31:*.z=00;31:*.Z=00;31:*.g
↳ z=00;31:*.bz2=00;31:*.tb2=00;31:*.tz2=00;31:*.tb
↳ z2=00;31:*.xz=00;31:*.avi=01;35:*.bmp=01;35:*.fl
↳ i=01;35:*.gif=01;35:*.jpg=01;35:*.jpeg=01;35:*.m
↳ ng=01;35:*.mov=01;35:*.mpg=01;35:*.pcx=01;35:*.p
↳ bm=01;35:*.pgm=01;35:*.png=01;35:*.ppm=01;35:*.t
↳ ga=01;35:*.tif=01;35:*.xbm=01;35:*.xpm=01;35:*.d
↳ l=01;35:*.gl=01;35:*.wmv=01;35:*.aiff=00;32:*.au
↳ =00;32:*.mid=00;32:*.mp3=00;32:*.ogg=00;32:*.voc
↳ =00;32:*.wav=00;32:
LD_LIBRARY_PATH=/opt/cray/pe/papi/5.6.0.3/lib64:/opt
↳ /cray/job/2.2.3-6.0.7.0.44.1_g6c4e934.ari/lib64
PE_FFTW_DEFAULT_TARGET_interlagos=interlagos
PE_LIBSCI_DEFAULT_VOLATILE_PRGENV=CRAY GNU INTEL
PE_PETSC_DEFAULT_GENCOMPILERS_INTEL_interlagos=16.0
PE_TPSL_64_DEFAULT_GENCOMPILERS_INTEL_mic_knl=16.0
PE_TPSL_DEFAULT_GENCOMPS_CRAY_x86_64=86
PE_TRILINOS_DEFAULT_GENCOMPILERS_GNU_x86_64=71 53 49
PE_TRILINOS_DEFAULT_GENCOMPILERS_INTEL_x86_64=160
CRAY_IAA_INFO_FILE=/tmp/cray_iaa_info.12841473
SINFO_FORMAT=%9P %5a %8s %10l %6c %6z %7D %10T %N
```

```
CRAY_RCA_POST_LINK_OPTS=-L/opt/cray/rca/2.2.18-6.0.7
↳ .0_33.3_g2aa4f39.ari/lib64
↳ -lrca
PE_LIBSCI_PKGCONFIG_VARIABLES=PE_LIBSCI_OMP_REQUIRES
↳ _@openmp@:PE_SCI_EXT_LIBPATH:PE_SCI_EXT_LIBNAME
PE_PETSC_DEFAULT_VOLATILE_PRGENV=CRAY CRAY64 GNU
↳ GNU64 INTEL INTEL64
PE_PKGCONFIG_LIBS=mpich:AtpSigHandler:cray-rca:libsc
↳ i_mpi:libsci
PE_TPSL_64_DEFAULT_GENCOMPILERS_GNU_sandybridge=7.1
↳ 5.3 4.9
PE_TPSL_64_DEFAULT_GENCOMPILERS_INTEL_haswell=16.0
PE_MPICH_FIXED_PRGENV=INTEL
XNLSPATH=/usr/share/X11/nls
FTN_X86_64=/opt/cray/pe/cce/8.7.3/cce/x86_64
PE_PETSC_DEFAULT_GENCOMPILERS_CRAY_mic_knl=8.6
PE_PETSC_DEFAULT_GENCOMPILERS_CRAY_x86_64=8.6
PE_PETSC_DEFAULT_GENCOMPILERS_INTEL_skylake=16.0
PE_PETSC_DEFAULT_GENCOMPS_GNU_interlagos=71 53 49
PE_PETSC_DEFAULT_GENCOMPS_GNU_sandybridge=71 53 49
PE_PETSC_DEFAULT_GENCOMPS_INTEL_interlagos=160
PE_PETSC_DEFAULT_GENCOMPS_INTEL_sandybridge=160
PE_TPSL_DEFAULT_GENCOMPS_GNU_haswell=71 53 49
CRAY_CXX_IPA_LIBS=/opt/cray/pe/cce/8.7.3/cce/x86_64/
↳ lib/libcray-c++-rts.a
MPICH_ABORT_ON_ERROR=1
PE_LIBSCI_DEFAULT_GENCOMPS_CRAY_x86_64=86
PE_PAPI_DEFAULT_PKGCONFIG_VARIABLES=PE_PAPI_ACCEL_LI
↳ BS@accelerator@
PE_PETSC_DEFAULT_GENCOMPILERS_CRAY_haswell=8.6
PE_PETSC_DEFAULT_GENCOMPS_GNU_mic_knl=53
PE_PETSC_DEFAULT_GENCOMPS_INTEL_mic_knl=160
PE_TPSL_64_DEFAULT_GENCOMPILERS_GNU_interlagos=7.1
↳ 5.3 4.9
PE_TPSL_64_DEFAULT_GENCOMPS_INTEL_sandybridge=160
MPICH_DIR=/opt/cray/pe/mpt/7.7.2/gni/mpich-cray/8.6
SRUN_DEBUG=3
HOSTTYPE=x86_64
ATP_POST_LINK_OPTS=-Wl,-L/opt/cray/pe/atp/2.1.2/libA
↳ pp/
PE_FFTW_DEFAULT_REQUIRED_PRODUCTS=PE_MPICH
PE_FFTW_DEFAULT_TARGET_sandybridge=sandybridge
PE_HDF5_PARALLEL_DEFAULT_REQUIRED_PRODUCTS=PE_MPICH
PE_NETCDF_HDF5PARALLEL_DEFAULT_REQUIRED_PRODUCTS=PE_
↳ HDF5_PARALLEL
PE_PETSC_DEFAULT_GENCOMPILERS_INTEL_sandybridge=16.0
PE_TPSL_64_DEFAULT_GENCOMPILERS_CRAY_haswell=8.6
PE_MPICH_FORTRAN_PKGCONFIG_LIBS=mpichf90
TMOUT=259200
PE_PETSC_DEFAULT_GENCOMPILERS_GNU_mic_knl=5.3
RCLOCAL_PRGENV=true
APPS=/apps/daint
FROM_HEADER=
CHPL_CG_CPP_LINES=1
OFFLOAD_INIT=on_start
PE_LIBSCI_DEFAULT_GENCOMPILERS_INTEL_x86_64=16.0
```

```

PE_LIBSCI_GENCOMPS_INTEL_x86_64=160
PE_PRODUCT_LIST=CRAYPE_HASWELL:CRAY_RCA:CRAY_ALPS:DV
  ↳ S:CRAY_XPMEM:CRAY_DMAPP:CRAY_PMI:CRAY_UGNI:CRAY_
  ↳ UDREG:CRAY_LIBSCI:CRAYPE:CRAY:PERFTOOLS:CRAYPAT
PE_TPSL_DEFAULT_GENCOMPILERS_CRAY_x86_64=8.6
PE_TPSL_DEFAULT_GENCOMPS_GNU_interlagos=71 53 49
ALPS_LLI_STATUS_OFFSET=1
PAGER=less
PE_MPICH_DEFAULT_GENCOMPS_PGI=153
PE_PETSC_DEFAULT_GENCOMPILERS_GNU_x86_64=7.1 5.3 4.9
PE_TPSL_DEFAULT_GENCOMPS_GNU_x86_skylake=71 61
CRAY_MPICH_ROOTDIR=/opt/cray/pe/mpt/7.7.2
ALPS_APP_PE=0
CSHEDIT=emacs
PE_LIBSCI_GENCOMPILERS_GNU_x86_64=7.1 6.1 5.1 4.9
PE_PETSC_DEFAULT_GENCOMPS_GNU_skylake=61
PE_PETSC_DEFAULT_GENCOMPS_INTEL_skylake=160
PE_TPSL_64_DEFAULT_GENCOMPILERS_INTEL_x86_64=16.0
PE_MPICH_GENCOMPILERS_CRAY=8.6
PE_MPICH_MODULE_NAME=cray-mpich
XDG_CONFIG_DIRS=/etc/xdg
CRAYPAT_ROOT=/opt/cray/pe/perftools/7.0.3
PE_LIBSCI_DEFAULT_GENCOMPILERS_CRAY_x86_64=8.6
PE_LIBSCI_GENCOMPS_CRAY_x86_64=86
PE_MPICH_DEFAULT_VOLATILE_PRGENV=CRAY GNU PGI
PE_MPICH_TARGET_VAR_nvidia20=-lcudart
PE_TPSL_64_DEFAULT_REQUIRED_PRODUCTS=PE_MPICH:PE_LIB
  ↳ SCI
PE_TPSL_DEFAULT_GENCOMPS_CRAY_haswell=86
PE_TPSL_DEFAULT_GENCOMPS_CRAY_sandybridge=86
MINICOM=-c on
LIBGL_DEBUG=quiet
CRAY_DMAPP_INCLUDE_OPTS=-I/opt/cray/dmapp/7.1.1-6.0.
  ↳ 7.0.34.3__g5a674e0.ari/include
  ↳ -I/opt/cray/gni-headers/5.0.12.0-6.0.7.0.24.1__g
  ↳ 3b1768f.ari/include
CRAY_LIBSCI_BASE_DIR=/opt/cray/pe/libsci/18.07.1
CRAY_LIBSCI_DIR=/opt/cray/pe/libsci/18.07.1
DVS_VERSION=0.9.0
NLSPATH=/opt/cray/pe/cce/8.7.3/cce/x86_64/share/nls/
  ↳ En/%N.cat
PE_LIBSCI_ACC_DEFAULT_VOLATILE_PKGCONFIG_PATH=/opt/c
  ↳ ray/pe/libsci_acc/18.07.1/@PRGENV@/@PE_LIBSCI_AC
  ↳ C_DEFAULT_GENCOMPS@/@PE_LIBSCI_ACC_DEFAULT_TARGE
  ↳ T@/lib/pkgconfig
PE_LIBSCI_PKGCONFIG_LIBS=libsci_mpi:libsci
PE_NETCDF_DEFAULT_GENCOMPS_GNU=
PE_PARALLEL_NETCDF_DEFAULT_GENCOMPS_GNU=51 49
PE_TPSL_64_DEFAULT_GENCOMPS_GNU_mic_knl=71 53
PE_TPSL_64_DEFAULT_GENCOMPS_GNU_x86_64=71 53 49
MODULE_VERSION=3.2.10.6

```

```

PAT_REPORT_PRUNE_NAME=_cray$mt_execute_,_cray$mt_sta
  ↳ rt_,_cray_hwpc_,f_cray_hwpc_,cstart_,_pat_,pat_
  ↳ region_,PAT_,OMP.slave_loop,slave_entry,_new_sla
  ↳ ve_entry,_thread_pool_slave_entry,THREAD_POOL_jo
  ↳ in_,_libc_start_main_,start_,_start_,start_thread
  ↳ ,_wrap_,UPC_ADIO_,_upc_,upc_,_caf_,_pgas_,sys_
  ↳ call_,_device_stub
PE_HDF5_DEFAULT_VOLATILE_PKGCONFIG_PATH=/opt/cray/pe
  ↳ /hdf5/1.10.2.0/@PRGENV@/@PE_HDF5_DEFAULT_GENCOMP
  ↳ S@/lib/pkgconfig
PE_PKGCONFIG_DEFAULT_PRODUCTS=PE_TRILINOS:PE_TPSL_64
  ↳ :PE_TPSL:PE_PETSC:PE_PARALLEL_NETCDF:PE_NETCDF_H
  ↳ DF5PARALLEL:PE_NETCDF:PE_MPICH:PE_LIBSCI_ACC:PE_
  ↳ LIBSCI:PE_HDF5_PARALLEL:PE_HDF5:PE_GA:PE_FFTW
PE_TPSL_DEFAULT_GENCOMPILERS_GNU_x86_64=7.1 5.3 4.9
PE_TPSL_DEFAULT_GENCOMPS_CRAY_interlagos=86
PE_MPICH_GENCOMPILERS_GNU=7.1 5.1 4.9
PMI_CRAY_NO_SMP_ORDER=0
CPU=x86_64
CSCS_CUSTOM_ENV=true
XTPE_NETWORK_TARGET=aries
ATP_IGNORE_SIGTERM=1
PE_FFTW_DEFAULT_TARGET_abudhabi=abudhabi
PE_MPICH_DEFAULT_DIR_PGI_DEFAULT64=64
PE_NETCDF_DEFAULT_GENCOMPILERS_GNU=7.1 6.1 5.3 4.9
PE_PARALLEL_NETCDF_DEFAULT_GENCOMPILERS_GNU=5.1 4.9
PE_PETSC_DEFAULT_GENCOMPS_CRAY_mic_knl=86
PE_TPSL_64_DEFAULT_GENCOMPILERS_GNU_x86_skylake=7.1
  ↳ 6.1
PE_TPSL_DEFAULT_GENCOMPILERS_GNU_haswell=7.1 5.3 4.9
_=/usr/bin/env
PMI_NO_FORK=1
JAVA_BINDIR=/usr/lib64/jvm/java/bin
QUEUE_SORT=-t,e,S
CRAY_CCE_SHARE=/opt/cray/pe/cce/8.7.3/cce/x86_64/sha
  ↳ re
CRAY_CXX_IPA_LIBS_AARCH64=/opt/cray/pe/cce/8.7.3/cce
  ↳ /aarch64/lib/libcray-c++-rts.a
PE_HDF5_PARALLEL_DEFAULT_FIXED_PRGENV=CRAY PGI INTEL
PE_HDF5_PARALLEL_DEFAULT_GENCOMPS_GNU=
PE_NETCDF_HDF5PARALLEL_DEFAULT_FIXED_PRGENV=CRAY PGI
  ↳ INTEL
PE_NETCDF_HDF5PARALLEL_DEFAULT_GENCOMPS_GNU=
PE_SMA_DEFAULT_DIR_CRAY_DEFAULT64=64
PE_TPSL_64_DEFAULT_GENCOMPILERS_CRAY_x86_skylake=8.6
CRAY_UDREG_POST_LINK_OPTS=-L/opt/cray/udreg/2.3.2-6.
  ↳ 0.7.0.33.18__g5196236.ari/lib64
PE_TPSL_64_DEFAULT_GENCOMPS_CRAY_sandybridge=86
PE_TPSL_64_DEFAULT_VOLATILE_PRGENV=CRAY CRAY64 GNU
  ↳ GNU64 INTEL INTEL64
PE_TPSL_DEFAULT_GENCOMPILERS_CRAY_mic_knl=8.6
PE_TPSL_DEFAULT_GENCOMPS_INTEL_interlagos=160
CRAYPE_VERSION=2.5.15
CRAY_ALPS_POST_LINK_OPTS=-L/opt/cray/alps/6.6.43-6.0
  ↳ .7.0.26.4__ga796da3.ari/lib64
PE_TPSL_DEFAULT_GENCOMPS_GNU_mic_knl=71 53

```



## Red-Blue Pebbling Revisited: Near Optimal Parallel Matrix Multiplication

```

PE_MPICH_VOLATILE_PRGENV=CRAY GNU PGI
JAVA_HOME=/usr/lib64/jvm/java
TARGETMODULES=craype-abudhabi:craype-abudhabi-cu:craype-accel-host:craype-accel-nvidia20:craype-accel-nvidia30:craype-accel-nvidia35:craype-barcelona:craype-broadwell:craype-haswell:craype-hugepages128K:craype-hugepages128M:craype-hugepages16M:craype-hugepages256M:craype-hugepages2M:craype-hugepages32M:craype-hugepages4M:craype-hugepages512M:craype-hugepages64M:craype-hugepages8M:craype-intel-knc:craype-interlagos:craype-interlagos-cu:craype-istanbul:craype-ivybridge:craype-mc12:craype-mc8:craype-mic-knl:craype-network-aries:craype-network-gemini:craype-network-infiniband:craype-network-none:craype-network-seastar:craype-sandybridge:craype-shanghai:craype-target-compute_node:craype-target-local_host:craype-target-native:craype-xeon:xtpe-barcelona:xtpe-interlagos:xtpe-interlagos-cu:xtpe-istanbul:xtpe-mc12:xtpe-mc8:xtpe-network-gemini:xtpe-network-seastar:xtpe-shanghai:xtpe-target-native:xtpe-xeon
INCLUDE_PATH_X86_64=/opt/cray/pe/cce/8.7.3/cce/x86_64/include/craylibs
PE_LIBSCI_DEFAULT_OMP_REQUIRES=
PE_MPICH_DEFAULT_GENCOMPS_CRAY=86
PE_PETSC_DEFAULT_GENCOMPILERS_GNU_sandybridge=7.1 5.3 4.9
PE_TPSSL_DEFAULT_GENCOMPILERS_INTEL_haswell=16.0
XALT_TRANSMISSION_STYLE=directdb
PE_LIBSCI_ACC_DEFAULT_NV_SUFFIX_nvidia20=nv20
PE_LIBSCI_MODULE_NAME=cray-libsci/18.07.1
PE_PETSC_DEFAULT_GENCOMPILERS_CRAY_skylake=8.6
PE_TPSSL_DEFAULT_GENCOMPILERS_CRAY_interlagos=8.6
PE_TPSSL_DEFAULT_GENCOMPILERS_GNU_mic_knl=7.1 5.3
LANG=en_US.UTF-8
PE_TPSSL_64_DEFAULT_GENCOMPS_GNU_x86_skylake=71 61
PE_INTEL_FIXED_PKGCONFIG_PATH=/opt/cray/pe/mpt/7.7.2/gni/mpich-intel/16.0/lib/pkgconfig
PYTHONSTARTUP=/etc/pythonstart
MODULEPATH=/apps/daint/modulefiles:/apps/daint/system/modulefiles:/apps/daint/UES/easybuild/modulefiles:/apps/common/UES/modulefiles:/apps/common/system/modulefiles:/opt/cray/pe/perftools/7.0.3/modulefiles:/opt/cray/pe/craype/2.5.15/modulefiles:/opt/cray/pe/modulefiles:/opt/cray/modulefiles:/opt/cray/pe/ari/modulefiles
PE_LIBSCI_GENCOMPILERS_CRAY_x86_64=8.6
PE_MPICH_NV_LIBS_nvidia20=-lcudart
PE_MPICH_VOLATILE_PKGCONFIG_PATH=/opt/cray/pe/mpt/7.7.2/gni/mpich-@PRGENV@/PE_MPICH_DIR_DEFAULT64/@PE_MPICH_GENCOMPS@/lib/pkgconfig
SDK_HOME=/usr/lib64/jvm/java
SHMEM_ABORT_ON_ERROR=1

```

```

CRAY_BINUTILS_ROOT_X86_64=/opt/cray/pe/cce/8.7.3/bin/
utils/x86_64/x86_64-pc-linux-gnu/./
CRAY_DMAPP_POST_LINK_OPTS=-L/opt/cray/dmapp/7.1.1-6.0.7.0_34.3_g5a674e0.ari/lib64
PE_FFTW_DEFAULT_TARGET_ivybridge=ivybridge
PE_FFTW_DEFAULT_TARGET_share=share
PE_FFTW_DEFAULT_TARGET_x86_skylake=x86_skylake
PE_PKG_CONFIG_PATH=/opt/cray/pe/cti/1.0.7/lib/pkgconfig:/opt/cray/pe/cti/1.0.6/lib/pkgconfig:/opt/cray/pe/cti/1.0.4/lib/pkgconfig
PE_TPSSL_64_DEFAULT_GENCOMPS_GNU_interlagos=71 53 49
PE_TPSSL_DEFAULT_GENCOMPILERS_INTEL_mic_knl=16.0
CRAY_RCA_INCLUDE_OPTS=-I/opt/cray/rca/2.2.18-6.0.7.0_33.3_g2aa4f39.ari/include
-I/opt/cray/krca/2.2.4-6.0.7.1_5.27_g8505b97.ari/include
-I/opt/cray/hss-devel/8.0.0/include
PAT_BUILD_PAPI_BASEDIR=/opt/cray/pe/papi/5.6.0.3
PE_LIBSCI_OMP_REQUIRES_openmp=_mp
PE_PETSC_DEFAULT_GENCOMPILERS_GNU_skylake=6.1
PE_TPSSL_DEFAULT_GENCOMPILERS_CRAY_x86_skylake=8.6
PE_TPSSL_64_DEFAULT_GENCOMPS_CRAY_mic_knl=86
PE_TPSSL_DEFAULT_GENCOMPILERS_INTEL_x86_64=16.0
CRAY_MPICH_DIR=/opt/cray/pe/mpt/7.7.2/gni/mpich-cray/8.6
PE_MPICH_CXX_PKGCONFIG_LIBS=mpichcxx
SQUEUE_FORMAT=%8i %8u %7a %14j %3t %9r %19S %10M %10L %5D %4C
CRAY_BINUTILS_BIN_AARCH64=/opt/cray/pe/cce/8.7.3/bin/
utils/cross/x86_64-aarch64/aarch64-linux-gnu/bin
PE_LIBSCI_ACC_DEFAULT_GENCOMPILERS_GNU_x86_64=4.9
PE_LIBSCI_DEFAULT_GENCOMPS_INTEL_x86_64=160
PE_MPICH_PKGCONFIG_VARIABLES=PE_MPICH_NV_LIBS_@accelerator@:PE_MPICH_ALTERNATE_LIBS_@multithreaded@:PE_MPICH_ALTERNATE_LIBS_@dpme
ENVIRONMENT=BATCH
APP2_STATE=7.0.3
CRAY_CC_VERSION=8.7.3
CRAY_PMI_POST_LINK_OPTS=-L/opt/cray/pe/pmi/5.0.14/lib64
PE_HDF5_DEFAULT_FIXED_PRGENV=CRAY PGI INTEL
PE_TPSSL_64_DEFAULT_GENCOMPILERS_CRAY_mic_knl=8.6
PE_TPSSL_DEFAULT_GENCOMPILERS_INTEL_x86_skylake=16.0
PE_TPSSL_DEFAULT_VOLATILE_PKGCONFIG_PATH=/opt/cray/pe/tpsl/18.06.1/@PRGENV@/PE_TPSSL_DEFAULT_GENCOMPS_@PE_TPSSL_DEFAULT_TARGET@/lib/pkgconfig
CRAY_MPICH2_VER=7.7.2
PE_MPICH_PKGCONFIG_LIBS=mpich
GPG_TTY=not a tty
PE_GA_DEFAULT_GENCOMPILERS_GNU=5.3 4.9
PE_LIBSCI_ACC_DEFAULT_GENCOMPS_GNU_x86_64=49
PE_LIBSCI_VOLATILE_PKGCONFIG_PATH=/opt/cray/pe/libsci/18.07.1/@PRGENV@/PE_LIBSCI_GENCOMPS_@PE_LIBSCI_TARGET@/lib/pkgconfig
PE_MPICH_ALTERNATE_LIBS_multithreaded=_mt
PE_NETCDF_DEFAULT_FIXED_PRGENV=CRAY PGI INTEL

```

```

PE_PARALLEL_NETCDF_DEFAULT_FIXED_PRGENV=CRAY PGI
↪ INTEL
SHLVL=4
JDK_HOME=/usr/lib64/jvm/java
QT_SYSTEM_DIR=/usr/share/desktop-data
CRAY_LIBSCI_VERSION=18.07.1
PE_HDF5_PARALLEL_DEFAULT_VOLATILE_PRGENV=GNU
PE_MPICH_TARGET_VAR_nvidia35=-lcudart
PE_NETCDF_HDF5PARALLEL_DEFAULT_VOLATILE_PRGENV=GNU
PE_PKGCONFIG_PRODUCTS_DEFAULT=PE_PAPI
PE_TPSSL_64_DEFAULT_GENCOMPS_GNU_haswell=71 53 49
OSTYPE=linux
LESS_ADVANCED_PREPROCESSOR=no
PE_TPSSL_DEFAULT_GENCOMPILERS_INTEL_interlagos=16.0
LINKER_X86_64=/opt/cray/pe/cce/8.7.3/binutils/x86_64
↪ /x86_64-pc-linux-gnu/bin/ld
PE_LIBSCI_ACC_DEFAULT_NV_SUFFIX_nvidia60=nv60
PE_MPICH_DEFAULT_VOLATILE_PKGCONFIG_PATH=/opt/cray/p
↪ e/mpt/7.7.2/gni/mpich-@PRGENV@PE_MPICH_DEFAULT_
↪ DIR_DEFAULT64@/PE_MPICH_DEFAULT_GENCOMPS@/lib/p
↪ kgconfig
PE_PETSC_DEFAULT_GENCOMPILERS_CRAY_interlagos=8.6
PE_TPSSL_DEFAULT_VOLATILE_PRGENV=CRAY CRAY64 GNU GNU64
↪ INTEL INTEL64
XCURSOR_THEME=DMZ
LS_OPTIONS=-N --color=none -T 0
CRAY_PMI_INCLUDE_OPTS=-I/opt/cray/pe/pmi/5.0.14/incl
↪ ude
PE_TPSSL_64_DEFAULT_GENCOMPS_CRAY_interlagos=86
PE_TPSSL_DEFAULT_GENCOMPS_INTEL_sandybridge=160
WINDOWMANAGER=
PRGENVMODULES=PrgEnv-cray:PrgEnv-gnu:PrgEnv-intel:Pr
↪ gEnv-pathscales:PrgEnv-pgi
CRAYPE_NETWORK_TARGET=aries
ATP_MRNET_COMM_PATH=/opt/cray/pe/atp/2.1.2/libexec/a
↪ tp_mrnet_commnodewrapper
CRAYLMD_LICENSE_FILE=/opt/cray/pe/cce/cce.lic
PE_TPSSL_DEFAULT_GENCOMPILERS_CRAY_haswell=8.6
PKG_CONFIG_PATH_DEFAULT=/opt/cray/pe/papi/5.6.0.2/li
↪ b64/pkgconfig
PE_MPICH_DIR_CRAY_DEFAULT64=64
PE_LEVEL=8.7
PE_PETSC_DEFAULT_GENCOMPILERS_GNU_haswell=7.1 5.3 4.9
PE_TPSSL_64_DEFAULT_GENCOMPILERS_GNU_mic_knl=7.1 5.3
PE_TPSSL_DEFAULT_GENCOMPILERS_GNU_interlagos=7.1 5.3
↪ 4.9
PE_TPSSL_DEFAULT_GENCOMPILERS_INTEL_sandybridge=16.0
MACHTYPE=x86_64-suse-linux
LESS=-M -I -R
G_FILENAME_ENCODING=@locale,UTF-8,ISO-8859-15,CP1252
CRAYLIBS_AARCH64=/opt/cray/pe/cce/8.7.3/cce/aarch64/
↪ lib
CRAYLIBS_X86_64=/opt/cray/pe/cce/8.7.3/cce/x86_64/lib
CRAY_GNI_HEADERS_INCLUDE_OPTS=-I/opt/cray/gni-header
↪ s/5.0.12.0-6.0.7.0_24.1__g3b1768f.ari/include

```

```

CRAY_LIBSCI_PREFIX_DIR=/opt/cray/pe/libsci/18.07.1/C
↪ RAY/8.6/x86_64
PE_HDF5_DEFAULT_GENCOMPS_GNU=
PE_MPICH_NV_LIBS=
PE_NETCDF_DEFAULT_REQUIRED_PRODUCTS=PE_HDF5
PE_TPSSL_64_DEFAULT_GENCOMPILERS_GNU_haswell=7.1 5.3
↪ 4.9
PE_TPSSL_64_DEFAULT_GENCOMPILERS_INTEL_sandybridge=16
↪ .0
PE_TPSSL_DEFAULT_GENCOMPS_GNU_x86_64=71 53 49
PE_TRILINOS_DEFAULT_REQUIRED_PRODUCTS=PE_MPICH:PE_HD
↪ F5_PARALLEL:PE_NETCDF_HDF5PARALLEL:PE_LIBSCI:PE_
↪ TPSL
PYTHONPATH=/apps/daint/UES/xalt/0.7.6/site:/apps/dai
↪ nt/UES/xalt/0.7.6/libexec
DMAPP_ABORT_ON_ERROR=1
PE_LIBSCI_OMP_REQUIRES=
PE_MPICH_DEFAULT_GENCOMPILERS_CRAY=8.6
PE_TRILINOS_DEFAULT_GENCOMPS_GNU_x86_64=71 53 49
PE_MPICH_GENCOMPS_CRAY=86
XDG_DATA_DIRS=/usr/share
TOOLMODULES=apprentice:apprentice2:atp:chapel:cray-l
↪ gdb:craypat:craypkg-gen:cray-snp launcher:ddt:gdb
↪ :iobuf:papi:perftools:perftools-lite:stat:totalv
↪ iew:xt-craypat:xt-lgdb:xt-papi:xt-totalview
DVS_INCLUDE_OPTS=-I/opt/cray/dvs/2.7.2.2.113-6.0.7.1
↪ .7.6__g1bbc03e/include
PE_LIBSCI_ACC_DEFAULT_GENCOMPILERS_CRAY_x86_64=8.5
PE_LIBSCI_ACC_DEFAULT_NV_SUFFIX_nvidia35=nv35
PE_LIBSCI_DEFAULT_REQUIRED_PRODUCTS=PE_MPICH
PE_MPICH_DEFAULT_FIXED_PRGENV=INTEL
PE_MPICH_DEFAULT_GENCOMPS_GNU=71 51 49
PE_TPSSL_64_DEFAULT_GENCOMPILERS_INTEL_interlagos=16.0
PE_TPSSL_DEFAULT_GENCOMPILERS_CRAY_sandybridge=8.6
MODULESHOME=/opt/cray/pe/modules/3.2.10.6
PE_GA_DEFAULT_FIXED_PRGENV=CRAY PGI INTEL
PE_LIBSCI_DEFAULT_VOLATILE_PKGCONFIG_PATH=/opt/cray/
↪ pe/libsci/18.07.1/@PRGENV@/PE_LIBSCI_DEFAULT_GE
↪ NCOMPS@/PE_LIBSCI_DEFAULT_TARGET@/lib/pkgconfig
PE_TPSSL_DEFAULT_GENCOMPILERS_GNU_sandybridge=7.1 5.3
↪ 4.9
LESSOPEN=lessopen.sh %s
CRAY_CXX_IPA_LIBS_X86_64=/opt/cray/pe/cce/8.7.3/cce/
↪ x86_64/lib/libcray-c++-rts.a
PE_MPICH_NV_LIBS_nvidia35=-lcudart
PE_PETSC_DEFAULT_VOLATILE_PKGCONFIG_PATH=/opt/cray/p
↪ e/petsc/3.8.4.0/complex/@PRGENV@/PE_PETSC_DEFAU
↪ LT_GENCOMPS@/PE_PETSC_DEFAULT_TARGET@/lib/pkgco
↪ nfig
PELOCAL_PRGENV=true
CRAY_NUM_COOKIES=2
CRAYPAT_OPTS_EXECUTABLE=sbin/pat-opts
CRAY_BINUTILS_BIN_X86_64=/opt/cray/pe/cce/8.7.3/binu
↪ tils/x86_64/bin
INCLUDE_PATH_AARCH64=/opt/cray/pe/cce/8.7.3/cce/aarc
↪ h64/include/craylibs

```

## Red-Blue Pebbling Revisited: Near Optimal Parallel Matrix Multiplication

```
LIBSCI_BASE_DIR=/opt/cray/pe/libsci/18.07.1
PE_TPSL_64_DEFAULT_GENCOMPS_INTEL_x86_64=160
CRAY_COOKIES=1768161280,1768226816
LIBSCI_VERSION=18.07.1
PE_LIBSCI_DEFAULT_PKGCONFIG_VARIABLES=PE_LIBSCI_DEFA
↳ ULT_OMP_REQUIRES=@openmp@:PE_SCI_EXT_LIBPATH:PE_
↳ SCI_EXT_LIBNAME
PE_MPICH_NV_LIBS_nvvidia60=-lcudart
PE_TPSL_64_DEFAULT_GENCOMPS_GNU_sandybridge=71 53 49
PE_TPSL_DEFAULT_GENCOMPS_INTEL_mic_knl=160
CRAY_PRE_COMPILE_OPTS=-hnetwork=aries
CRAY_ALPS_INCLUDE_OPTS=-I/opt/cray/alps/6.6.43-6.0.7
↳ .0_26.4__ga796da3.ari/include
PE_FFTW_DEFAULT_TARGET_broadwell=broadwell
PE_LIBSCI_GENCOMPILERS_INTEL_x86_64=16.0
PE_PGI_DEFAULT_FIXED_PKGCONFIG_PATH=/opt/cray/pe/par
↳ allel-netcdf/1.8.1.3/PGI/15.3/lib/pkgconfig:/opt
↳ /cray/pe/netcdf-hdf5parallel/4.6.1.2/PGI/17.10/1
↳ ib/pkgconfig:/opt/cray/pe/netcdf/4.6.1.2/PGI/17.
↳ 10/lib/pkgconfig:/opt/cray/pe/hdf5-parallel/1.10
↳ .2.0/PGI/17.10/lib/pkgconfig:/opt/cray/pe/hdf5/1
↳ .10.2.0/PGI/17.10/lib/pkgconfig:/opt/cray/pe/ga/
↳ 5.3.0.8/PGI/17.10/lib/pkgconfig
PE_TPSL_64_DEFAULT_GENCOMPILERS_GNU_x86_64=7.1 5.3
↳ 4.9
CRAY_CPU_TARGET=haswell
CRAY_UGNI_INCLUDE_OPTS=-I/opt/cray/ugni/6.0.14.0-6.0
↳ .7.0_23.1__gea11d3d.ari/include
CRAY_XPMEM_INCLUDE_OPTS=-I/opt/cray/xpmem/2.2.15-6.0
↳ .7.1.5.10__g7549d06.ari/include
PE_LIBSCI_REQUIRED_PRODUCTS=PE_MPICH
PE_MPICH_DEFAULT_GENCOMPILERS_PGI=15.3
PE_PAPI_DEFAULT_ACCELL_FAMILY_LIBS=
PE_TPSL_64_DEFAULT_GENCOMPS_CRAY_x86_64=86
craype_already_loaded=0
PE_MPICH_GENCOMPS_PGI=153
PE_LIBSCI_DEFAULT_GENCOMPILERS_GNU_x86_64=7.1 6.1 5.1
↳ 4.9
PE_LIBSCI_GENCOMPS_GNU_x86_64=71 61 51 49
PE_TPSL_DEFAULT_GENCOMPS_INTEL_haswell=160
LESSCLOSE=lessclose.sh %s %s
ATP_HOME=/opt/cray/pe/atp/2.1.2
PE_FFTW_DEFAULT_TARGET_x86_64=x86_64
PE_PETSC_DEFAULT_GENCOMPILERS_INTEL_x86_64=16.0
G_BROKEN_FILENAMES=1
CC_X86_64=/opt/cray/pe/cce/8.7.3/cce/x86_64
```

```
CRAY_LD_LIBRARY_PATH=/opt/cray/pe/mpt/7.7.2/gni/mpic
↳ h-cray/8.6/lib:/opt/cray/pe/perftools/7.0.3/lib6
↳ 4:/opt/cray/rca/2.2.18-6.0.7.0_33.3__g2aa4f39.ari
↳ i/lib64:/opt/cray/alps/6.6.43-6.0.7.0_26.4__ga79
↳ 6da3.ari/lib64:/opt/cray/xpmem/2.2.15-6.0.7.1_5.
↳ 10__g7549d06.ari/lib64:/opt/cray/dmapp/7.1.1-6.0
↳ .7.0_34.3__g5a674e0.ari/lib64:/opt/cray/pe/pmi/5
↳ .0.14/lib64:/opt/cray/ugni/6.0.14.0-6.0.7.0_23.1
↳ __gea11d3d.ari/lib64:/opt/cray/udreg/2.3.2-6.0.7
↳ .0_33.18__g5196236.ari/lib64:/opt/cray/pe/libsci
↳ /18.07.1/CRAY/8.6/x86_64/lib:/opt/cray/pe/cce/8.
↳ 7.3/cce/x86_64/lib
PE_FFTW_DEFAULT_TARGET_haswell=haswell
PE_GA_DEFAULT_GENCOMPS_GNU=53 49
PE_GA_DEFAULT_VOLATILE_PKGCONFIG_PATH=/opt/cray/pe/g
↳ a/5.3.0.8/@PRGENV@/@PE_GA_DEFAULT_GENCOMPS@/lib/
↳ pkgconfig
+ lsb_release -a
PE_INTEL_DEFAULT_FIXED_PKGCONFIG_PATH=/opt/cray/pe/p
↳ arallel-netcdf/1.8.1.3/INTEL/16.0/lib/pkgconfig:
↳ /opt/cray/pe/netcdf-hdf5parallel/4.6.1.2/INTEL/1
↳ 6.0/lib/pkgconfig:/opt/cray/pe/netcdf/4.6.1.2/IN
↳ TEL/16.0/lib/pkgconfig:/opt/cray/pe/mpt/7.7.2/gn
↳ i/mpich-intel/16.0/lib/pkgconfig:/opt/cray/pe/hd
↳ f5-parallel/1.10.2.0/INTEL/16.0/lib/pkgconfig:/o
↳ pt/cray/pe/hdf5/1.10.2.0/INTEL/16.0/lib/pkgconfi
↳ g:/opt/cray/pe/ga/5.3.0.8/INTEL/18.0/lib/pkgconf
↳ ig
PE_PAPI_DEFAULT_ACCEL_LIBS=
PE_PETSC_DEFAULT_GENCOMPILERS_GNU_interlagos=7.1 5.3
↳ 4.9
PE_PETSC_DEFAULT_GENCOMPILERS_INTEL_haswell=16.0
PE_SMA_DEFAULT_DIR_PGI_DEFAULT64=64
PE_TPSL_64_DEFAULT_GENCOMPILERS_INTEL_x86_skylake=16
↳ .0
COLORTERM=1
JAVA_ROOT=/usr/lib64/jvm/java
PE_MPICH_DEFAULT_DIR_CRAY_DEFAULT64=64
PE_PETSC_DEFAULT_GENCOMPS_CRAY_haswell=86
PE_PETSC_DEFAULT_GENCOMPS_GNU_x86_64=71 53 49
PE_PETSC_DEFAULT_GENCOMPS_INTEL_x86_64=160
BASH_FUNC_module%=() { eval
↳ `'/opt/cray/pe/modules/3.2.10.6/bin/modulecmd
↳ bash $*`
}
LSB Version: n/a
Distributor ID: SUSE
Description: SUSE Linux Enterprise Server 12 SP3
Release: 12.3
Codename: n/a
+ uname -a
Linux nid00008 4.4.103-6.38.4.0.153-cray_ari_c #1 SMP
↳ Thu Nov 1 16:05:05 UTC 2018 (gef8fef) x86_64
↳ x86_64 x86_64 GNU/Linux
+ lscpu
Architecture: x86_64
```



```

CPU op-mode(s):      32-bit, 64-bit
Byte Order:          Little Endian
CPU(s):              72
On-line CPU(s) list: 0-71
Thread(s) per core:  2
Core(s) per socket:  18
Socket(s):           2
NUMA node(s):        2
Vendor ID:            GenuineIntel
CPU family:           6
Model:                79
Model name:           Intel(R) Xeon(R) CPU E5-2695 v4
    ↳ @ 2.10GHz
Stepping:             1
CPU MHz:              2101.000
CPU max MHz:          2101.0000
CPU min MHz:          1200.0000
BogoMIPS:             4200.30
Virtualization:       VT-x
L1d cache:            32K
L1i cache:            32K
L2 cache:             256K
L3 cache:             46080K
NUMA node0 CPU(s):    0-17,36-53
NUMA node1 CPU(s):    18-35,54-71
Flags:                fpu vme de pse tsc msr pae mce
    ↳ cx8 apic sep mtrr pge mca cmov pat pse36 clflush
    ↳ dts acpi mmx fxsr sse sse2 ss ht tm pbe syscall nx
    ↳ pdpe1gb rdtscp lm constant_tsc arch_perfmon pebs
    ↳ bts rep_good nopl xtopology nonstop_tsc
    ↳ aperfmperf eagerfpu pni pclmulqdq dtes64 monitor
    ↳ ds_cpl vmx smx est tm2 ssse3 sdbg fma cx16 xtpr
    ↳ pdcm pcid dca sse4_1 sse4_2 x2apic movbe popcnt
    ↳ tsc_deadline_timer aes xsave avx f16c rdrand
    ↳ lahf_lm abm 3dnowprefetch ida arat epb
    ↳ invpcid_single pln pts dtherm intel_pt spec_ctrl
    ↳ kaiser tpr_shadow vnmi flexpriority ept vpid
    ↳ fsgsbase tsc_adjust bmi1 hle avx2 smep bmi2 erms
    ↳ invpcid rtm cqm rdseed adx smap xsaveopt cqm_llc
    ↳ cqm_occup_llc
+ cat /proc/meminfo
MemTotal:             65860732 kB
MemFree:               62682788 kB
MemAvailable:         62453220 kB
Buffers:               13916 kB
Cached:                936516 kB
SwapCached:            0 kB
Active:                347604 kB
Inactive:              889248 kB
Active(anon):          316824 kB
Inactive(anon):        802956 kB
Active(file):          30780 kB
Inactive(file):        86292 kB
Unevictable:           8184 kB
Mlocked:               8184 kB
SwapTotal:             0 kB

```

```

SwapFree:              0 kB
Dirty:                 0 kB
Writeback:             0 kB
AnonPages:             295320 kB
Mapped:                81864 kB
Shmem:                 831936 kB
Slab:                  246364 kB
SReclaimable:          24980 kB
SUnreclaim:            221384 kB
KernelStack:           14272 kB
PageTables:            4900 kB
NFS_Unstable:           0 kB
Bounce:                0 kB
WritebackTmp:          0 kB
CommitLimit:           32930364 kB
Committed_AS:          1112368 kB
VmallocTotal:          34359738367 kB
VmallocUsed:            0 kB
VmallocChunk:           0 kB
HardwareCorrupted:     0 kB
HugePages_Total:       0
HugePages_Free:         0
HugePages_Rsvd:         0
HugePages_Surp:         0
Hugepagesize:           2048 kB
DirectMap4k:            11000 kB
DirectMap2M:            1974272 kB
DirectMap1G:            67108864 kB
+ inxi -F -c0
+ lsblk -a
NAME MAJ:MIN RM SIZE RO TYPE MOUNTPOINT
loop0  7:0    0 18.6M 0 loop
    ↳ /var/opt/cray/imps-distribution/squash/mounts/p0
loop1  7:1    0 128K 0 loop /var/opt/cray/imps-dis-
    ↳ tribution/squash/mounts/global
loop2  7:2    0 2.5G 1 loop /.rootfs_lower_ro
loop3  7:3    0 1.6G 1 loop
    ↳ /var/opt/cray/imps-image-binding/diags/squash_mo
    ↳ unts/squashfs_03SZzt_mount_point
loop4  7:4    0      1 loop
loop5  7:5    0      0 loop
loop6  7:6    0      0 loop
loop7  7:7    0      0 loop
+ lsscsi -s
+ module list
++ /opt/cray/pe/modules/3.2.10.6/bin/modulecmd bash
    ↳ list
Currently Loaded Modulefiles:
  1) modules/3.2.10.6
  2) gcc/6.2.0
  3) craype-broadwell
  4) craype-network-aries
  5) craype/2.5.15
  6) cray-mpich/7.7.2
  7) slurm/17.11.12.cscs-1
  8) xalt/daint-2016.11

```

- 9) daint-mc
- 10) udreg/2.3.2-6.0.7.1\_5.13\_\_g5196236.ari
- 11) ugni/6.0.14-6.0.7.1\_3.13\_\_gea11d3d.ari
- 12) pmi/5.0.14
- 13) dmapp/7.1.1-6.0.7.1\_5.45\_\_g5a674e0.ari
- 14) gni-headers/5.0.12-6.0.7.1\_3.11\_\_g3b1768f.ari
- 15) xpmem/2.2.15-6.0.7.1\_5.11\_\_g7549d06.ari
- 16) job/2.2.3-6.0.7.1\_5.43\_\_g6c4e934.ari
- 17) dvs/2.7\_2.2.118-6.0.7.1\_10.2\_\_g58b37a2
- 18) alps/6.6.43-6.0.7.1\_5.45\_\_ga796da32.ari
- 19) rca/2.2.18-6.0.7.1\_5.47\_\_g2aa4f39.ari
- 20) atp/2.1.2
- 21) perftools-base/7.0.3
- 22) PrgEnv-gnu/6.0.4
- 23) CMake/3.12.4
- 24) intel/18.0.2.199

## ARTIFACT EVALUATION

*Verification and validation studies:* Computation correctness: The correctness of all computations were cross-validated between all implementations and presence of unit tests for COSMA library.

Communication volume: For algorithms for which we verified the decomposition and communication strategy statically (CARMA and COSMA), we validated the communication volume measured by the mpiP profiler by our communication models.

We also evaluated total runtime of COSMA breakdown to communication and computation routines, using the semiprof profiler (<https://github.com/bcumming/semiprof>)

*Accuracy and precision of timings:* We measured the time with the standard std::chrono c++ library + MPI\_Barrier. The precision of chrono was set to milliseconds, which is a much higher precision than necessary for our experiments.

Each of the performance experiments were performed six times. The first run was always discarded due to the setup overhead. For the next five runs, we computed arithmetic mean and 95% confidence intervals. The limited number of runs was dictated by large scale of the experiments (up to 512 compute nodes).

*Used manufactured solutions or spectral properties:* N/A

*Quantified the sensitivity of results to initial conditions and/or parameters of the computational environment:* N/A

*Controls, statistics, or other steps taken to make the measurements and analyses robust to variability and unknowns in the system.* We took 3285 performance measurements in total for 6 different scenarios, ranging from 100 to 18432 cores (3 to 512 compute nodes). For communication volume, we gathered 372 data points (up to 64 compute nodes). Both the matrix sizes and number of cores were chosen to represent a wide spectrum of both performance-friendly and real-world use cases.

The repetitive runs were performed in different time to compensate for different machine occupancy and node allocations.